

Serie 4 - Solution

Exercise 1 (Debugging). FIXME : Nico...

Exercise 2 (Profiling). Producing the gmon.out file is done with :

```
gcc -pg poisson.c -o poisson
```

We get then :

```

index % time    self  children    called    name
                                <spontaneous>
[1]    100.0    5.30    5.36
              5.36    0.00    1019/1019    main [1]
              0.00    0.00     2/2          write_to_file [2]
              0.00    0.00     1/1          second [3]
              0.00    0.00     1/1          colormap [4]
-----
              5.36    0.00    1019/1019    main [1]
[2]    50.3    5.36    0.00    1019    write_to_file [2]
-----
              0.00    0.00     2/2          main [1]
[3]    0.0    0.00    0.00     2          second [3]
-----
              0.00    0.00     1/1          main [1]
[4]    0.0    0.00    0.00     1          colormap [4]
-----

```

Most of the time is spent in writing the files. The file is written in ASCII so it is not parallelizable. One need to modify into a binary format.

Exercise 3 (icc, gcc, optimization, vectorization and other cool stuff). Results are done on a two-CPU Intel Haswell E5-2630v3 running at 2.4 GHz with DDR4. $N = 1024$, $eps = 0.005$

optimization level	t_{icc} [s]	t_{gcc} [s]
-O0	58.92	57.32
-O1	9.08	9.39
-O2	3.34	9.12
-O3	3.36	8.67
-O3 -xHost (-ftree-vectorize)	2.74	8.67

A simple way to measure the execution time is using `second()` :

Exercise 4 (Parallelization with OpenMP). `int main() {`

```

    int    i, j, k;
    float **u;
    float **uo;
    float **f;
    float  h;
    float  l2;

```

```

double t1;

t1 = second();

h = 1. / (N+1);

// Allocation of arrays of pointers
u = (float**) malloc((N+1)*sizeof(float*));
uo = (float**) malloc((N+1)*sizeof(float*));
f = (float**) malloc((N+1)*sizeof(float*));

// allocation of the rows
for(i=0; i<N+1 ;i++) {
    u [i] = (float*) malloc((N+1)*sizeof(float));
    uo[i] = (float*) malloc((N+1)*sizeof(float));
    f [i] = (float*) malloc((N+1)*sizeof(float));
}

// initialization of u0 and f
#pragma omp parallel
{
    #pragma omp for private(i,j)
    for(i = 0; i < N+1; i++) {
        for(j = 0; j < N+1; j++) {
            uo[i][j] = 0;
            u[i][j] = 0;
            f [i][j] = -2.*100. * M_PI * M_PI * sin(10.*M_PI*i*h) * sin(10.*M_PI*j*h);
        }
    }
}

k=0;
do {
    l2 = 0.;

    #pragma omp parallel
    {
        #pragma omp for reduction(+ : l2) private(i,j)
        for(i = 1; i < N; i++) {
            for(j = 1; j < N ;j++) {
                // computation of the new step
                u[i][j] = 0.25 * ( uo[i-1][j] + uo[i+1][j] + uo[i][j-1] + uo[i][j+1] - f[i][j]*h*h);

                // L2 norm
                l2 += (uo[i][j] - u[i][j])*(uo[i][j] - u[i][j]);
            }
        }
    }
}

```

```
}

    // copy new grid in old grid
#pragma omp parallel
{
    #pragma omp for private(i,j)
        for(i = 0; i < N+1; i++){
            for(j = 0; j < N+1; j++){
                uo[i][j] = u[i][j];
            }
        }
}

    // outputs
    printf("l2=%.5f (k=%d)\n", sqrt(l2), k);

//    write_to_file(N+1, u, k, -1., 1.);
    k++;
} while(l2 > eps);

colormap(N+1);

// deallocation of the rows
for(i=0; i<N+1 ;i++) {
    free(u [i]);
    free(uo[i]);
    free(f [i]);
}

// deallocate the pointers
free(u);
free(uo);
free(f);

printf("Execution time = %f [s] \n", (second()-t1));
return 0;
}
```

Results are done on a two-CPU Intel Haswell, 16 cores (32 with Hyperthreading) E5-2630v3 running at 2.4 GHz with DDR4. $N = 1024$, $eps = 0.005$

OMP_NUM_THREADS	time [s]	speedup
1	2.85	1
2	1.49	1.9
4	0.73	3.9
8	0.43	6.6
16	0.73	3.9
32	0.84	3.3

One can observe that the speedup is close to linear up to 8 cores running. After this limit, we observe a speeddown. This is due to the limit of the memory bandwidth (north bridge). Conclusion : this application runs at best with 8 cores on this machine.

