

Performance and Performance modeling

What is HPC?

What is HPC?

It's High Performance Computing...


... High Performance means minimize Time To Solution

Why?

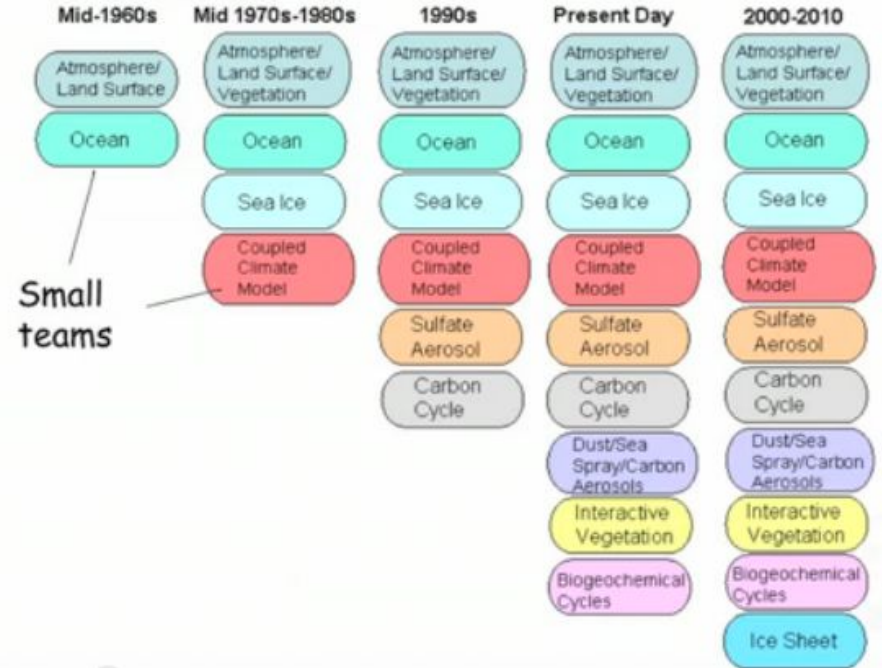
Why More Performance?

FAR TO GO

The Blue Brain Project has steadily increased the scale of its cortical simulations through the use of cutting-edge supercomputers and ever-increasing memory resources. But the full-scale simulation called for in the proposed Human Brain Project (red) would require resources roughly 100,000 times larger still.



Timeline of Climate Model Development



HPC First Principles:

Latency and Throughput

Latency: time to complete a operation

Throughput: how many operations per time

Water hose analogy:

- Latency is how long it takes to water to go through the pipe
- Throughput is the amount of water the pipe is outputting

Remarks:

- Cutting the pipe in half halves latency but throughput remains unchanged
- Adding a pipe does not change the latency but increases the throughput



HPC First Principles:

Latency and Throughput

- Easily confused (they refer to speed)
- Often contradict each other
- **Relate to each other with Little's Law**

$$L = \lambda W$$

L = average number of items in the queuing system

W = average waiting time in the system for an item

λ = average number of items arriving per unit time


Which one is the most important?



HPC First Principles: Latency and Throughput

In general, latency reduction hits physical limits

Top Airplane Speeds and Their Dates of Record, from Wright to Now




*For its overall records, FAI recognizes only manned, air-breathing jet aircraft, not rocket-powered airplanes like the Bell X-1 that Yeager flew.
Sources: Fédération Aéronautique Internationale (FAI), National Air and Space Museum, Concorde History, Top Speed, and Boeing.
Produced by Eli Dourado and Michael Kotrous, July 2016.

MERCATUS CENTER
George Mason University

Global air travel increase between 1974-2014

Data: World Bank. More: Bitsofscience.org. Hash: [#TryNotToFly](https://twitter.com/Bitsofscience)



Series : Air transport, passengers carried

Source: World Development Indicators

Created on: 02/24/2016

Planes do not fly faster but there are much more in the sky...

Latency and Throughput for CPUs

Latency (in seconds or cycles):

how long it takes before the next dependent operation can start

dependant operations performance is limited by latency

Throughput (in Instruction Per Cycles or Per Seconds): number of independent operations per time unit

independent operations performance is limited by throughput

Little's Law

parallelism = latency * throughput

Or

latency = parallelism/throughput

Or

throughput = parallelism/latency

i.e. latency is covered with parallelism



HPC First Principles: Memory Latency and Throughput

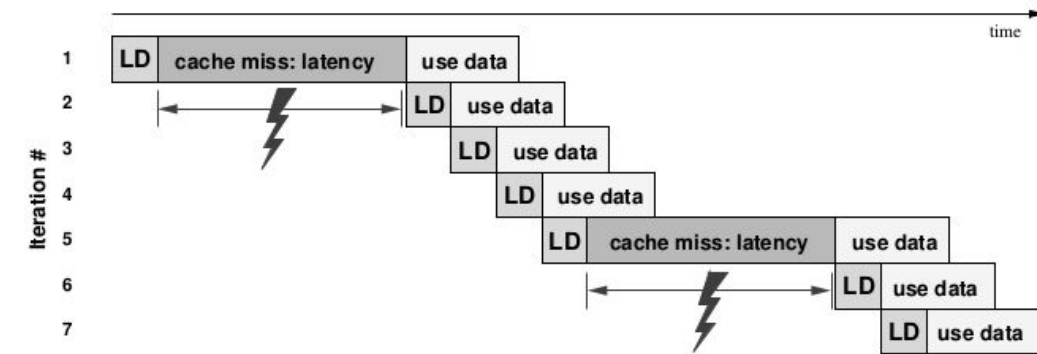
Memory bandwidth goes up (nice!) but latency does not go down (not nice).

SPEED VS. LATENCY AS MEMORY TECHNOLOGY HAS MATURED (INDUSTRY STANDARDS)				
TECHNOLOGY	MODULE SPEED (MT/s)	CLOCK CYCLE TIME (ns)	CAS LATENCY (CL)	TRUE LATENCY (ns)
SDR	100	8.00	3	24.00
SDR	133	7.50	3	22.50
DDR	335	6.00	2.5	15.00
DDR	400	5.00	3	15.00
DDR2	667	3.00	5	15.00
DDR2	800	2.50	6	15.00
DDR3	1333	1.50	9	13.50
DDR3	1600	1.25	11	13.75
DDR4	1866	1.07	13	13.93
DDR4	2133	0.94	15	14.06
DDR4	2400	0.83	17	14.17
DDR4	2666	0.75	18	13.50


<http://www.crucial.com/usa/en/memory-performance-speed-latency>

HPC First Principles: Memory Latency and Throughput

Memory Prefetchers: increasing throughput by maximizing locality



Latency bound



Bandwidth bound

CPU needs to know where to fetch data:

- Contiguous accesses will maximize throughput
- Non-contiguous accesses are latency bound

Maths Operations Latency and Throughput


	Broadwell		KNL		Kepler GPU	
	Lat (C)	TP (IPC)	Lat (C)	TP (IPC)	Lat (C)	TP (IPC)
Add	3	1	2	2	9	?
Multiply	3	2	7	2	9	?
Fused MAdd	5	2	6	2	9	32 (?)
Division	10-14	0.05	32	0.031	141	?
Sqrt	10-23	0.2	38	0.063	181	?
SinCos	52-124	?	50-250	0.006	18	?
Atan	97-147	?	125-265	0.001	?	?
Log	92	?	190	0.005	22	?

As a reference, a DRAM memory access is about 200 cycles (memory wall).

CPUs are designed to perform multiply-adds, but that's about it...


use as many FMAs as possible.

HPC First Principles: CPU Latency and Throughput




CPU's frequency has hit a limit more than 10 years ago but performance still goes up thanks to **Parallel Performance**


CPU Peak FP Performance



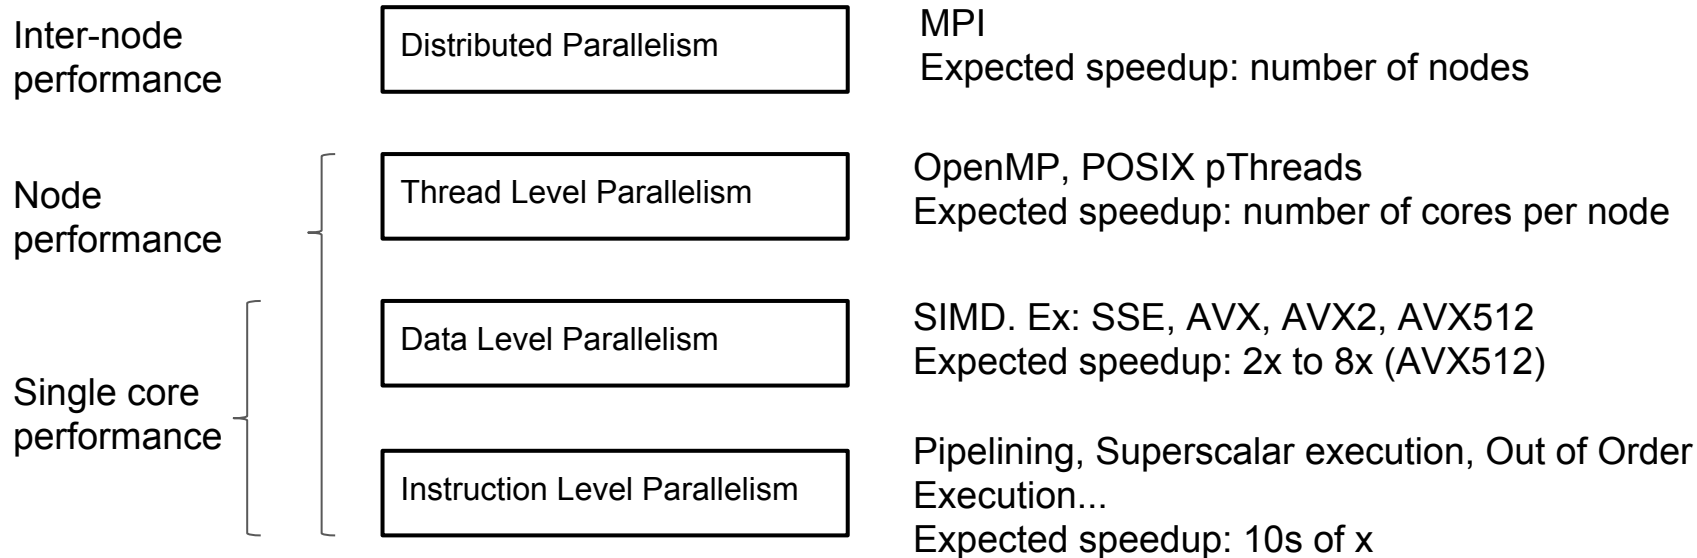
CPU Peak FP Performance - The Ninja Gap



CPU Peak FP Performance - The Ninja Gap



Different levels of parallelism



Performance is based on the duplication of resource and is harnessed by parallelism

Instruction Level Parallelism

Pipelining


Superscalar architecture

Out-Of-Order Execution

Speculative execution



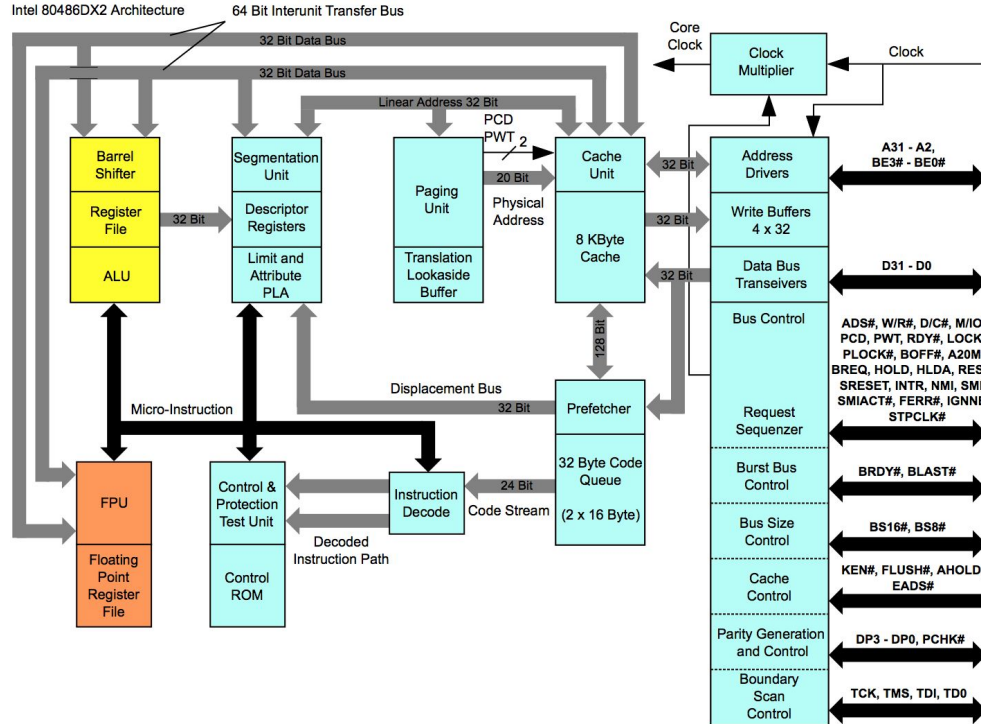
MELTDOWN




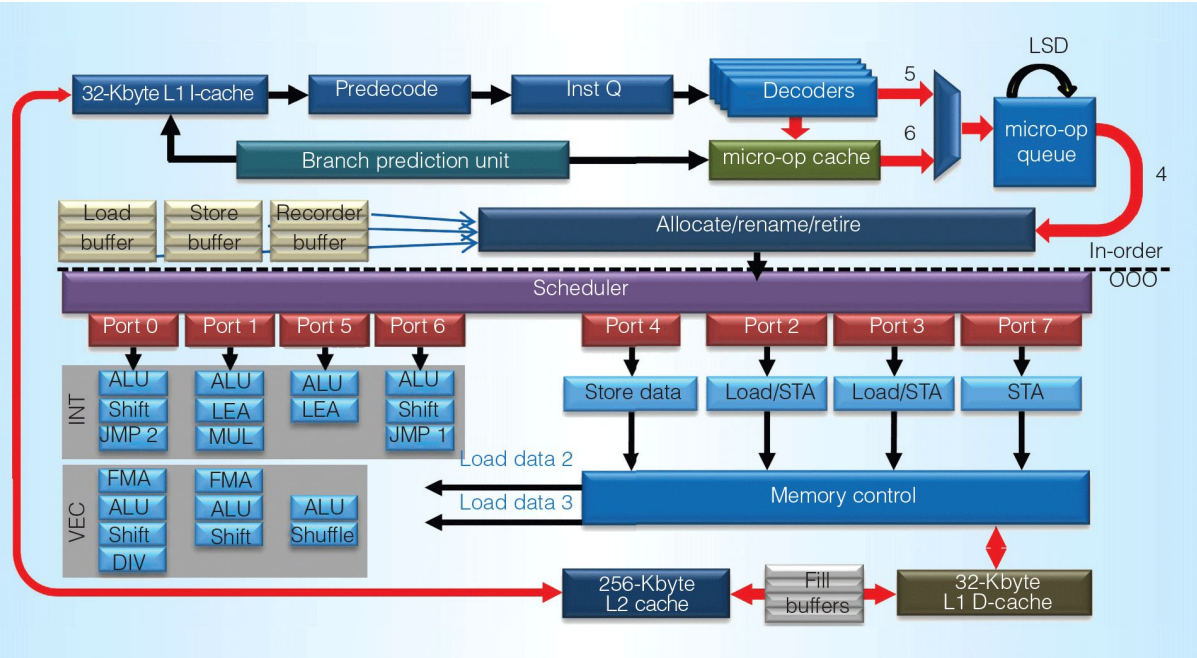
SPECTRE

The goal is to maximize CPU utilization and avoid stalls

Vintage CPU Architecture: Intel's 486 DX2

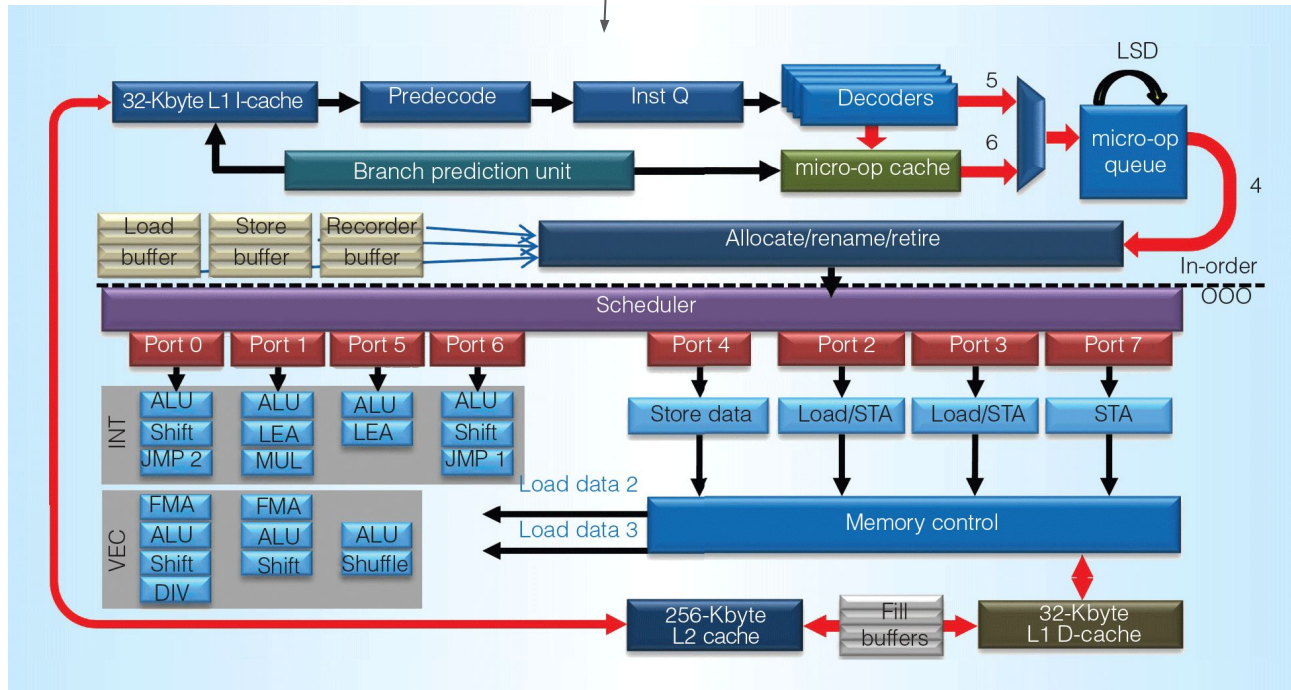


Modern CPU Architecture: Intel's Skylake




Bottom-up: mapping to the architecture

$$I_{\omega,1/2}^* = \frac{(1 - \epsilon^2) \kappa_o \omega}{i \sqrt{\epsilon}} \ln \frac{\frac{1-\epsilon}{1+\epsilon} x - i \frac{1+\epsilon}{1-\epsilon} y + 2i \sqrt{\epsilon} \sqrt{\omega^2 + \frac{x^2}{(1+\epsilon)^2} + \frac{y^2}{(1-\epsilon)^2}}}{x - iy + 2i \omega \sqrt{\epsilon}}.$$



Non-pipelined Processor



- 16 cycles to execute 4 instructions, **instruction latency is 4 cycles**
- **Throughput is $\frac{1}{4}$ Instructions Per Cycles**

Pipelining



Pipelining



Ford's assembly line

Pipelining


Ford's assembly line

Ford had been trying to increase his factories' productivity for years. The workers (...) arranged the parts in a row on the floor, put the under-construction auto on skids and dragged it down the line as they worked. **Ford broke the Model T's assembly into 84 discrete steps and trained each of his workers to do just one.**

The most significant piece of Ford's efficiency crusade was the assembly line.

<https://www.history.com/this-day-in-history/fords-assembly-line-starts-rolling>

Pipelining




- 16 cycles to execute 4 instructions broken down into 4 “stages”
- Instruction latency is **STILL** 4 cycles
- **Throughput** is $4/7$ ($\sim 1/2$) IPC, **almost 2X** compared to non-pipelined

Pipelining

Like the Ford's assembly line, instructions are broken down in many small steps (stages)

- Increased IPC through increased parallelism
- Smaller stages means increased frequency which unlocked the frequency era


But there's a price to pay: deep pipelines can easily stall




Pipelining

Stalls (bubbles) happen when the pipeline cannot advance properly

- Most probable cause is instruction dependency
- The CPU is waiting for a resource (e.g. read/write in the memory)



Superscalar Architecture



- Multiple pipelines to increase Instructions Per Cycles
- Can be spoiled by data, control, structural hazards and Multi-cycle instructions.


Out-Of-Order Execution

In-Order Execution: first instruction **in** is first instruction executed

Out-Of-Order Execution: first instruction **ready** is first instruction executed

- Operations are reordered
- Operations without dependencies are executed when the execution engines are ready
- Results are immediately available when needed (if the prediction was correct)

Out-Of-Order execution

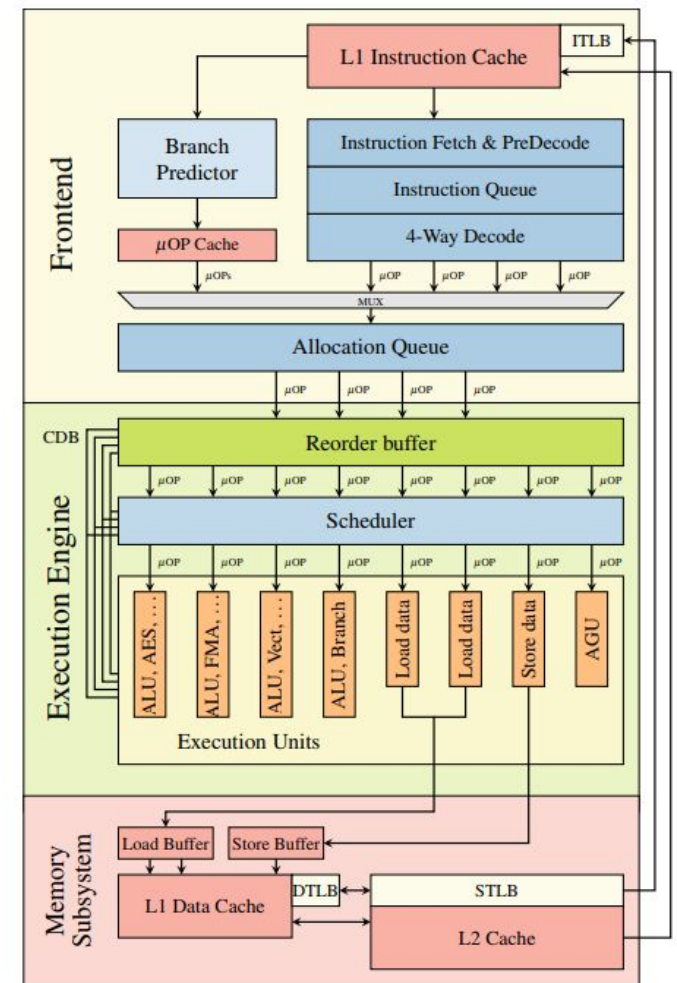
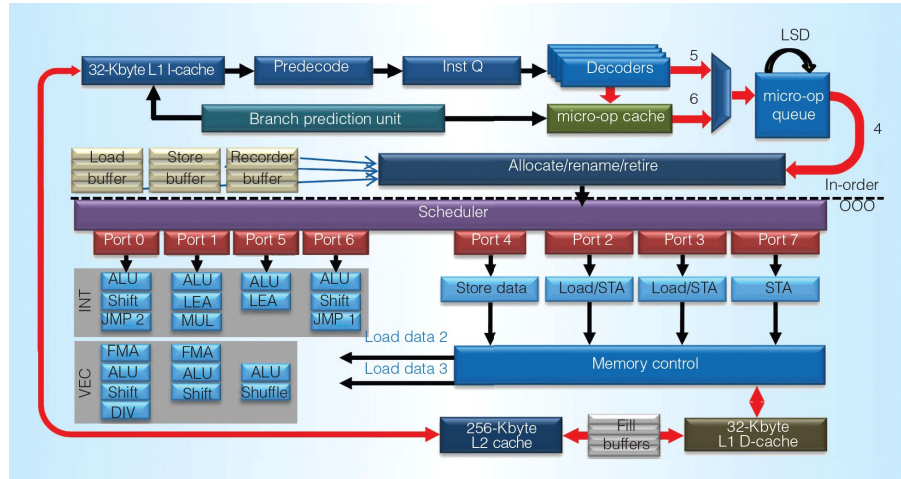


Modern CPU architecture

CPU Execution Engines:


- Pipelined
- Out-Of-Order execution
- Superscalar

ILP keeps the pipelines full **if parallelism is extracted**



Speculative Execution

Speculative execution: tentative **execution despite dependencies**




Very interesting read: <https://meltdownattack.com/meltdown.pdf>


Data Level Parallelism (SIMD)

SIMD: Same Instruction Multiple Data

Scalar operation



SIMD operation




- Throughput is multiplied by the vector size
- AVX: 256 bits, AVX512: 512 bits, GPUs: 2048 bits

HPC Data Structures: SoA vs AoS

```
struct {  
    float a[N];  
    float b[N];  
    float c[N];  
} SoA;
```

SOA




```
struct {  
    float a, b, c;  
} AoS[N];
```

AoS




HPC Data Structures: SoA vs AoS



- Information needs to be shuffled to and from the vector registers before and after the vector operations
- Compilers will have a very hard time optimizing/vectorizing the code
- Cache “unfriendly”

HPC Data Structures: SoA vs AoS



- 1-to-1 correspondence between the cache lines and the registers
- No shuffle/gather/scatter needed

DLP: Compilation and Assembly

Compilation flags (for gcc):

Standard Optimization flags:

- -O0: default, no optimization
- -O(1): enables optimization to reduce time to solution (for instance autovectorization/DLP)
- -O2: even more aggressive optimization

Standard Debug flags:

- -g0: no debug information
- -g1: minimal debug information
- -g: default debug information
- -g3: maximal debug information

Special flag: -Ofast

- sets -O3
- Sets **--fast-math**: extra performance boost you need but use with care
 - breaks strict IEEE-754 (floating point) compliance
 - Reorder code to improve ILP
 - enables fast approximation of transcendental functions (sqrt, div...)

DLP: Compilation and Assembly

Check the assembly for vector instructions

- Compilers options (-O)
- Misalignments,
- Standards were not designed for HPC,
- ...

Example:

vfmadd213pd

v: vector instruction

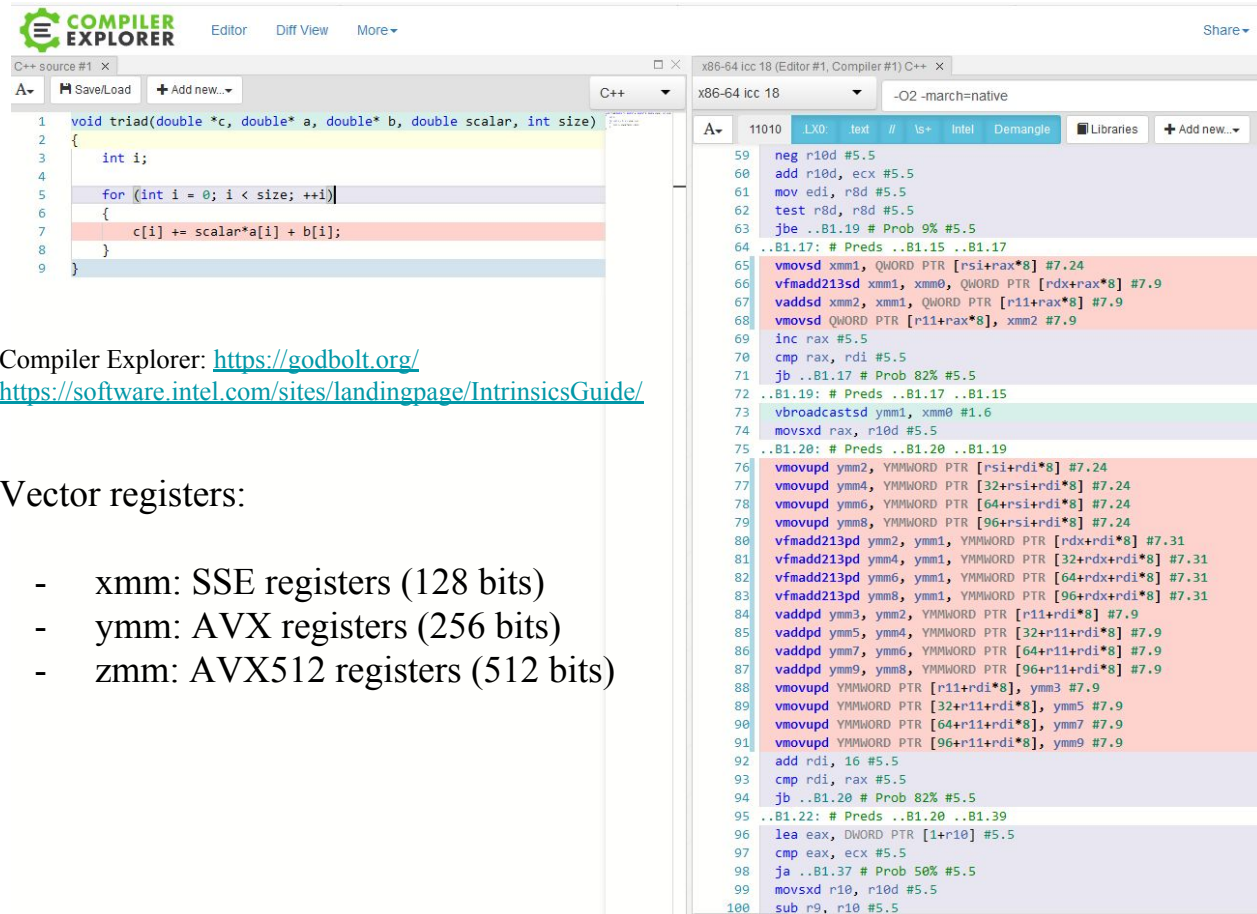
p: packed

d: double

vfmadd213ss

first s: scalar

second s: float (single)



The screenshot displays the Compiler Explorer interface. On the left, the C++ source code for a function named `triad` is shown. The function takes a pointer to a double array `a`, a scalar `b`, and an integer `size`. It iterates over the array, adding `scalar*a[i] + b[i]` to `c[i]`. On the right, the generated x86-64 assembly is displayed. The assembly includes instructions for setting up registers, loading data, and performing vector operations using SSE and AVX registers (xmm, ymm, zmm). The assembly is color-coded to match the source code, with red highlighting for the main loop and blue for initialization and cleanup.

Compiler Explorer: <https://godbolt.org/>
<https://software.intel.com/sites/landingpage/IntrinsicsGuide/>

Vector registers:

- xmm: SSE registers (128 bits)
- ymm: AVX registers (256 bits)
- zmm: AVX512 registers (512 bits)

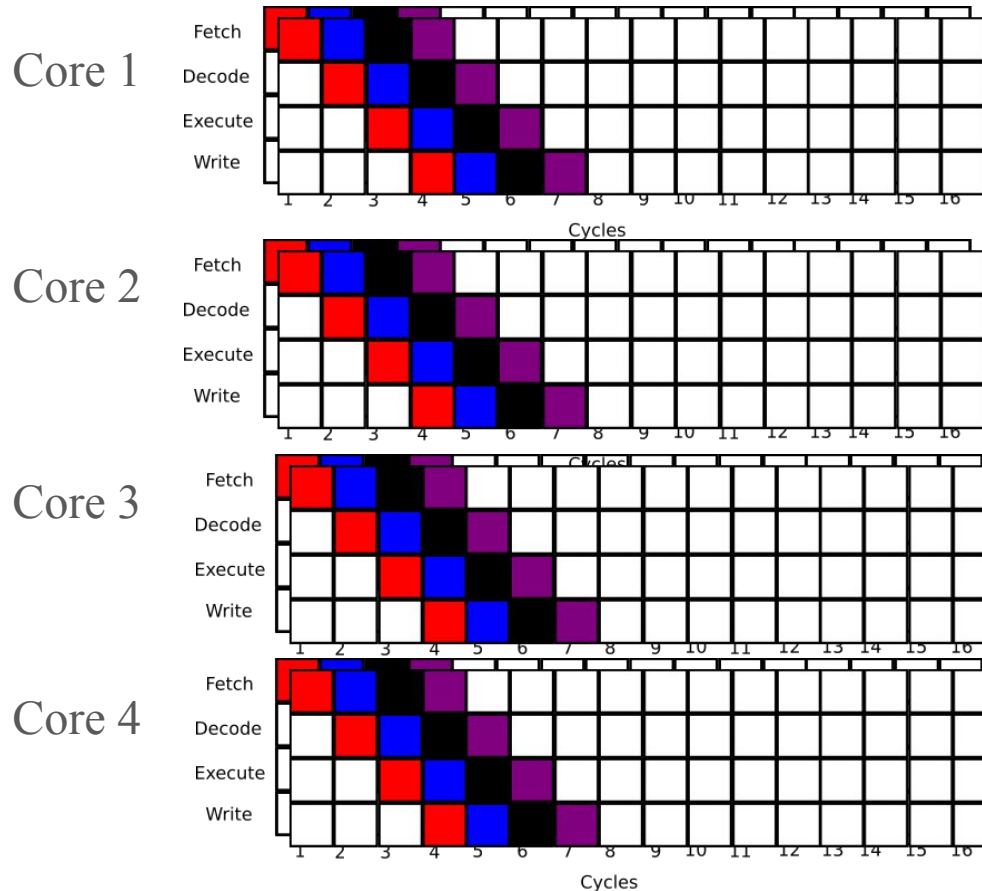
Thread Level Parallelism (multi-many core)

Use of multiple **concurrent** threads of execution
that are **inherently parallel**

- Increase throughput of multithreaded codes by covering the latencies
- Extremely effective in some cases (GPUs)

However...

- (just) resource (cores) duplication
- Diminishing returns
- Concurrent programming is very difficult
- Current hardware is not really designed to support TLP (e.g. caches...)



HPC Take Home Message - Node Performance

The goal of HPC is to increase mathematical throughput:

- Latency is NOT going down therefore throughput is increased
- Throughput is going up IF parallelism is increased
- Avoid pipeline stalls by having data “close” to the CPU

HPC kernel optimization focus on **extracting parallelism** and **maximizing data locality**

Assessing Performance: The Roofline Model

- There are many different types hardware (CPU, GPUs...)
- There are many different codes (algorithms, kernels...)

Is there a unified model to assess software performance?

Answer: the Roofline Model

- Roofline: an insightful visual performance model for multicore architectures, Williams, S. and Waterman, A. and Patterson, D., Communication to ACM, 2009
- A view of the parallel computing landscape, K. Asanovic, R. Bodik, J. Demmel, T. Keaveny, K. Keutzer, J. Kubiatowicz, N. Morgan, D. Patterson, K. Sen, J. Wawrzynek, D. Wessel, and K. Yelick, Communication to ACM, 2009

Assessing Performance: The Roofline Model

Software abstraction:

- Kernels can be represented by:
 - Number of mathematical operations (Flops)
 - Number of Data transfers (B)

Hardware abstraction:


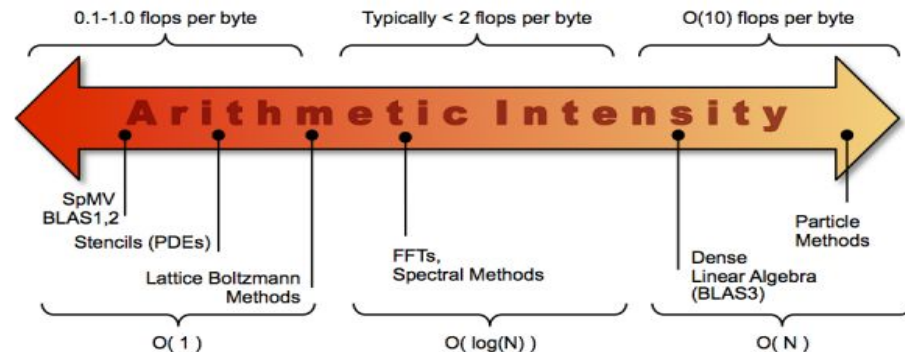
- Time To Solution is inversely proportional to:
 - Mathematical throughput (Flops/s)
 - DRAM throughput (B/s)

Assessing Performance: The Roofline Model

AI = flops/DRAM accesses


$$\text{Perf (flops)} = \min \left\{ \begin{array}{l} \text{Peak FP performance} \\ \text{Peak BW} * \text{AI} \end{array} \right\}$$

- Peak BW is measured with stream (<https://www.cs.virginia.edu/stream/>)
- Peak FP is measured with dgemm (BLAS 3)



Assessing Performance: The Roofline Model


Optimized Roofline:



The attainable system performance is the maximal performance that can be reached by an application:

$$\text{Attainable perf Gflops/s} = \min \left\{ \begin{array}{l} \text{Peak performance Gflops/s} \\ \text{Peak Memory Bandwidth} \times \text{Arithmetic Intensity} \end{array} \right.$$

Cache Aware Roofline:



•Total volume communication across all memory hierarchies relatively to the core:

$$\text{Attainable perf Gflops/s} = \min \left\{ \begin{array}{l} \text{Peak performance Gflops/s} \\ \text{Bandwidths to Core} \times \text{Arithmetic Intensity} \end{array} \right.$$

- Bandwidths to Core
- Bandwidth from L1 to Core
 - Bandwidth from L2 to Core
 - Bandwidth from L3 to Core
 - Bandwidth from DRAM to Core

CPU Theoretical peak FP performance

(Theoretical) **Peak FP performance** (Gflops/s) =
 Number of FP ports * (Superscalar architecture)
 flops/cycles * (e.g. 2 for FMAs)
 Vector size * (DLP)
 Frequency * (in Ghz)
 Number of cores (TLP)

Example: Skylake 2.5 Ghz, 28 cores (Platinum 8180)

Peak perf = 2 (FMA ports) * 2 (FMA flops) * 8 (DP) * 2.5
 = 32 flops DP/cycle * 2.5 Ghz
 = 80 Gflops/s DP per core
 = 2240 Gflops/s DP per socket (CPU)

Nina Gap Peak perf = $1 * 1 * 1 * 2.5 = 2.5$ Gflops/s, $\sim 1/1000$ th of the full peak

Assessing Performance: The Roofline Model

Advantages:

- Caps the performance
- Allows visual goals for optimization
- Easily used everywhere

Disadvantages:

- Latencies need to be covered
- Oversimplification
- Node-only
- “Vanilla” version is cache-oblivious
- OI can be hard to compute (e.g. Read For Ownership)

Assessing Performance: examples

Setup:

- Intel(R) Xeon(R) CPU E5-2680 v3 @ 2.50GHz
- DDR3 2.133 Ghz

Peak performance (single core): $2 \times 2 \times 4 \times 2.50 = 40$ Gflops/s, 34.5 measured

Peak Bandwidth (single core): ~ 20 GB/s (est.), 13.9 measured

Ridge Point: 2 Flops/B, 2.48 measured


- Triad:

$$c[i] = \text{scalar} * a[i] + b[i]$$

Arithmetic Intensity: $2 \text{ flops} / 3 * 8 \text{ Bytes} = 2/24 = 1/12 = 0.083$ Flops/B

Performance: $13.9 * 0.083 = 1.15$ Gflops

Assessing Performance: examples



Assessing Performance: examples

DGEMM: Double precision GEneralized Matrix Multiplication (BLAS 3)


$$C[i, j] += \text{sum}(A[i, k] * B[k, j]) \quad i, j, k = 1..N$$

Arithmetic Intensity:


$$2N^3 \text{ Flops}, 3N^2 * 8 \text{ DRAM accesses} = \text{N/12 flops/B}$$

Performance:

$$34.5 \text{ Gflops if } N > 2.35 * 3$$



Assessing Performance: examples



Assessing Performance: Another Go

Flops/cycle can be computed by:

$\text{Flops/cycle} = \text{number of FMA ports} * 2 (\text{FMA}) * \text{vec length}$

Example: Intel(R) Xeon(R) CPU E5-2680 v3 @ 2.50GHz

$\text{Flops/cycle} = 2 * 2 * 4 = 16 \text{ Flops/cycle Double Precision}$

Use RDTSCP on x86_64:

```
static inline unsigned long long cycles()
{
    unsigned long long u;
    asm volatile ("rdtscp; \
        shlq $32, %%rdx; \
        orq %%rdx, %%rax; \
        movq %%rax, %0" \
        : "=q" (u) : : "%rax", "%rdx", "rcx");
    return u;
}
```

And derive performance by counting flops in the code.

Triad example: $2 (\text{FMA}) * 4 (\text{vector size}) = 8 \text{ flops in 3 cycles}$

Performance: $8/3 \sim 2.6 \text{ Flops/cycle}$, i.e. 6.5 Gflops/s

```
- callq gettimeofday@plt
vxorpd %xmm0,%xmm0,%xmm0
vxorpd %xmm1,%xmm1,%xmm1
lea 0x148(%rsp),%rax
vcvttsi (%rax),%xmm0,%xmm0
vcvttsi -0x8(%rax),%xmm1,%xmm1
vmadd 0x25a3(%rip),%xmm0,%xmm1 # 4038d0 <__dso_handle+0x608>
vmovsd %xmm1,0xae0(%rsp,%r13,8)
vmovup 0x2501(%rip),%ymm1 # 403840 <__dso_handle+0x578>
xor %r8,%r8
792: vmovup 0x4caba500(,%r8,8),%ymm0
vmadd 0x26860300(,%r8,8),%ymm1,%ymm0
vmovnt %ymm0,0x606100(,%r8,8)
add $0x4,%r8
cmp $0x4c4b400,%r8
2.34 jnb 792
vzerou
lea 0x150(%rsp),%rdi
lea 0x60(%rsp),%rsi
- callq gettimeofday@plt
vxorpd %xmm0,%xmm0,%xmm0
vxorpd %xmm1,%xmm1,%xmm1
lea 0x158(%rsp),%rax
vcvttsi (%rax),%xmm0,%xmm0
vcvttsi -0x8(%rax),%xmm1,%xmm1
vmadd 0x252a(%rip),%xmm0,%xmm1 # 4038d0 <__dso_handle+0x608>
lea 0xae0(%rsp,%r13,8),%rdx
vsbds (%rdx),%xmm1,%xmm2
vmovsd %xmm2,(%rdx)
inc %r13
cmp $0x64,%r13
t jl 5c0
mov 0x68(%rsp),%r12
mov $0x4035b0,%edi
mov 0x70(%rsp),%r13
mov 0x78(%rsp),%r14
mov 0x80(%rsp),%r15
mov 0x88(%rsp),%rbx
rdtscp
ret 0
```

```
1 second sleep, number of cycles = 2500292456
Function Best Rate MB/s Avg time Min time Max time
Copy: inf 0.000000 0.000000 0.000001
Scale: inf 0.000000 0.000000 0.000000
Add: inf 0.000000 0.000000 0.000001
Triad: 21061.0 0.091392 0.091164 0.098817
cycles : 2.873238
```

Floating Point Representation

All real numbers can be approximated by floating point numbers represented by

$$(-1)^s d.dd... d \times \beta^e$$

where:

- $d.dd... d$ is called the mantissa and has p digits,
 $d.dd... d = (d_0 + d_1\beta + d_2\beta^2 + d_3\beta^3 + \dots d_{p-1}\beta^{p-1})$
- β is the base
- e is the exponent
- s is the sign

Any real number can be represented by a linear combination of

$$0.5, 0.25, 0.125, 0.0625, \dots, (1/2^{p-1})$$

IEEE 754 defines two different floating point representation, single (32bits) and double precisions (64bits) in base two:

Precision	Sign	Exponent	Mantissa
Single precision	1	8	23 (+1)
Double precision	1	11	52 (+1)

Floating Point Representation

Base Convert: IEEE 754 Floating Point

Decimal

-0.1

32 bit – float

Decimal (exact)

-0.100000001490116119384765625

Binary

1 01111011 10011001100110011001101

Hexadecimal

BDCCCCD

64 bit – double

Decimal (exact)

-0.10000000000000000055511151231257827021181583404541015625

Binary

1 01111111011 1001100110011001100110011001100110011001100110011010

Hexadecimal

BFB999999999999A

Floating Point Representation

Compute:

$$f(x, y) = (333.75 - x^2)y^6 + x^2(11x^2y^2 - 121y^4 - 2) + 5.5y^8 + x/(2y)$$

For

$$x = 77617 \text{ and } y = 33096 \text{ (exactly represented in FP32)}$$

$$\text{FP32: } f(x, y) = 1.1726$$

$$\text{FP64: } f(x, y) = 1.17260394005318$$

$$\text{FP128: } f(x, y) = 1.17260394005318631$$

Floating Point Representation

Compute:

$$f(x, y) = (333.75 - x^2)y^6 + x^2(11x^2y^2 - 121y^4 - 2) + 5.5y^8 + x/(2y)$$

For

$$x = 77617 \text{ and } y = 33096 \text{ (exactly represented in FP32)}$$

$$\text{FP32: } f(x, y) = 1.1726$$

$$\text{FP64: } f(x, y) = 1.17260394005318$$

$$\text{FP128: } f(x, y) = 1.17260394005318631$$

The correct answer is:

$$f(x, y) = -0.82739605994682136814116509547981629$$

computational precision != result accuracy

Floating Point Arithmetic

\mathbb{R} is a field:

1- Closure:

$a+b$ and $a.b$ are in \mathbb{R}

2- Commutative laws:

$a + b = b + a$, $a.b = b.a$

3- Associative laws:

$(a + b) + c = a + (b + c)$

$(a.b).c = a.(b.c)$

4- distributive laws:

$a.(b + c) = a.b + a.c$

5- there are real numbers 0 and 1 such that

$a + 0 = 0 + a = a$ and $a.1 = 1.a = a$

6- for any real number a , there exists $-a$ such that

$a + (-a) = (-a) + a$

and if $a \neq 0$, there exists a^{-1} such that

$a.a^{-1} = a^{-1}.a = 1$

7- cancellation: if $a.b = a.c$ and $a \neq 0$ then $b = c$

It follows that:

8- $(a + b) - b = a$ and $a.(b/a) = b$

Floating Point Arithmetic

FP is **NOT** a field:

1- Closure:

$a+b$ and $a.b$ are in FP

2- Commutative laws:

$a + b = b + a$, $a.b = b.a$

3- ~~Associative laws:~~

~~$(a + b) + c = a + (b + c)$~~

~~$(a.b).c = a.(b.c)$~~

4- ~~distributive laws:~~

~~$a.(b + c) = a.b + a.c$~~

5- there are real numbers 0 and 1 such that

$a + 0 = 0 + a = a$ and $a.1 = 1.a = a$

6- for any real number a , there exists $-a$ such that

$a + (-a) = (-a) + a$

and if $a \neq 0$, there exists a^{-1} such that

$a.a^{-1} = a^{-1}.a = 1$

7- ~~cancellation: if $a.b = a.c$ and $a \neq 0$ then $b = c$~~

It follows that:

8- $(a + b) - b = a$ and ~~$a.(b/a) = b$~~

HPC brutal facts

HPC is about minimizing Time To Solution by **maximizing the throughput using parallelism**, not reducing latency

Without HPC techniques, **your code won't run faster on supercomputers** than on your workstation

CPUs are very good at doing **fused multiply-adds** but that's about it

HPC is about knowing **how your software AND the hardware best work together** to get maximum performance

HPC is about **hacking** your way around the language standard

You'll need to have a **look at the assembly**

Bibliography

Computer Architecture, A Quantitative Approach

Patterson, 2011

Introduction to High Performance Computing for Scientists and Engineers

Georg Hager, Gerhard Wellin, CRC Press

Roofline: An Insightful Visual Performance Model for Multicore Architectures

Samuel Williams, Andrew Waterman, and David Patterson

Floating-Point Computation

Pat H. Sterbenz

Exercices

<https://c4science.ch/source/phpc-2019/>

Directory: Serie03