

Parallel and High Performance Computing

Lecture 2: Development and execution in an HPC environment

March 1, 2021

Outline

Today exercise session

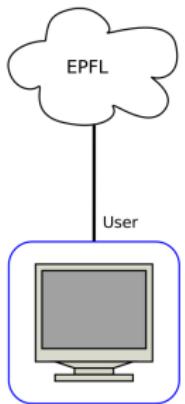
- What is a cluster.
- How to connect to a cluster (any remote machine).
- How to compile codes and use libraries.
- How to submit job on a cluster.
- How to use GIT

Getting your accounts

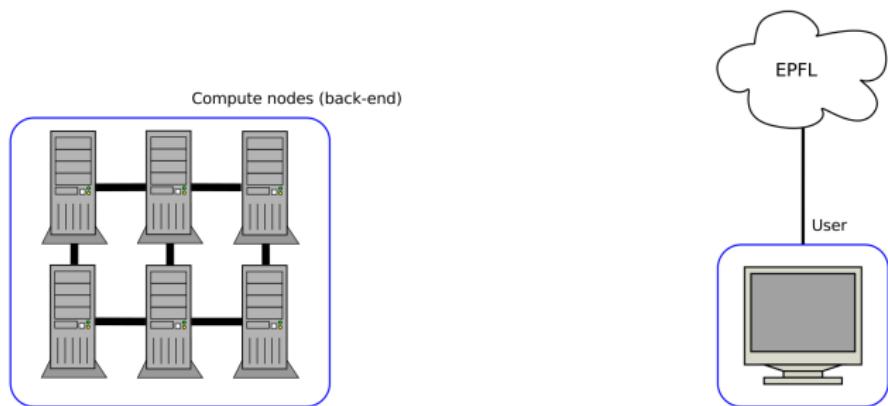
groups.epfl.ch

If you are not in the group **phpc-2021** let use now

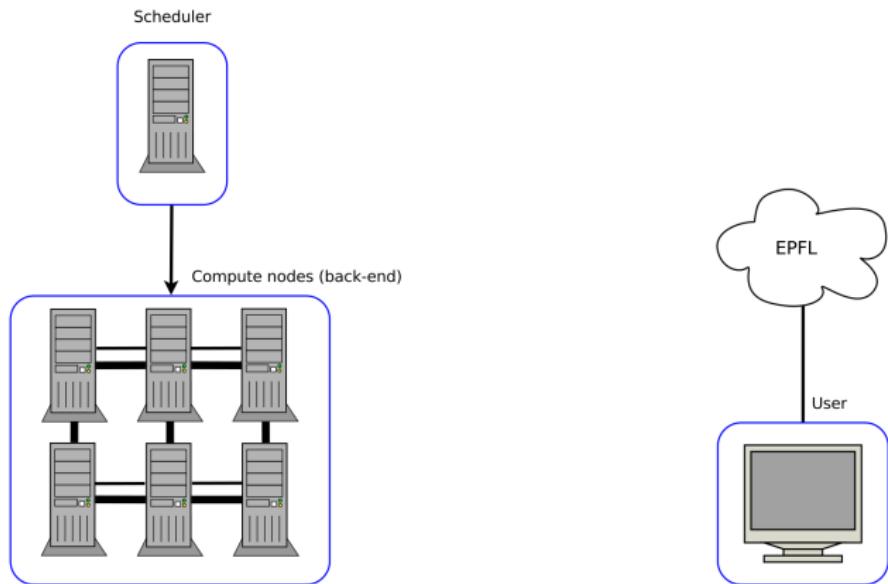
What is a cluster?



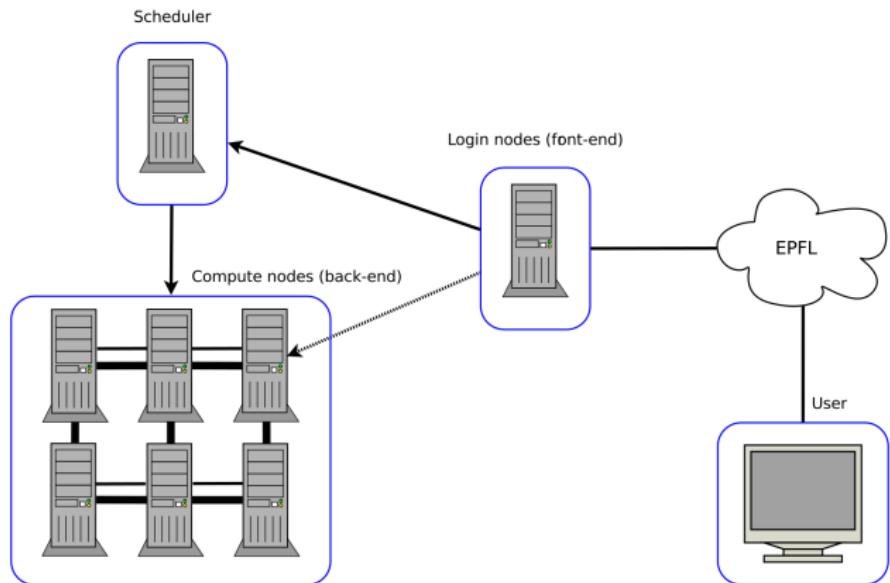
What is a cluster?



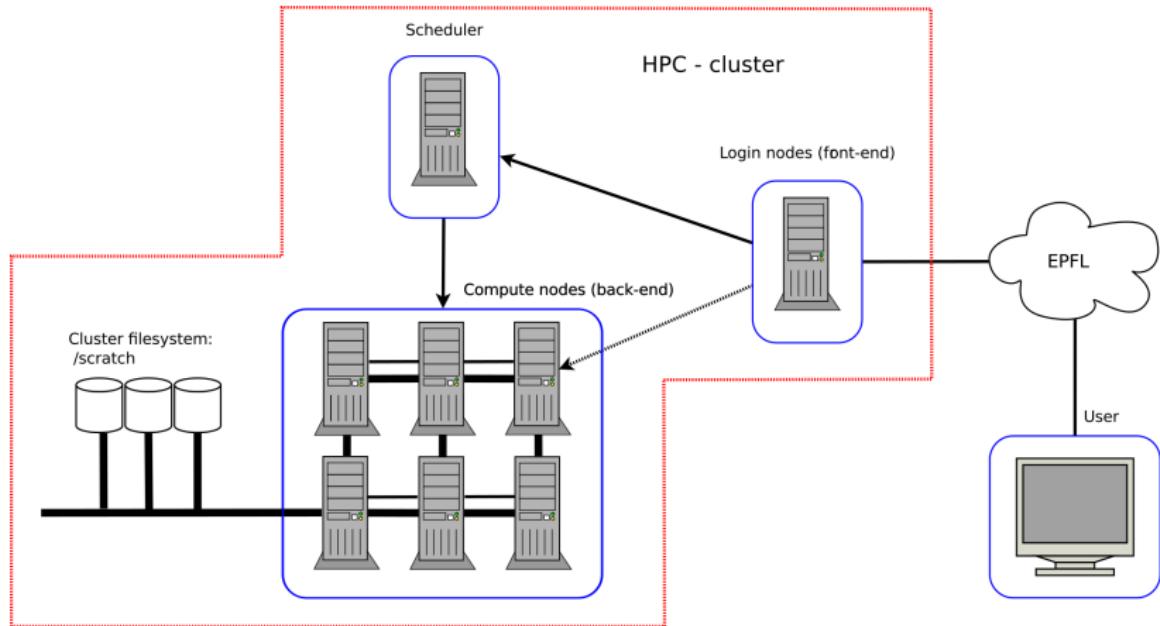
What is a cluster?



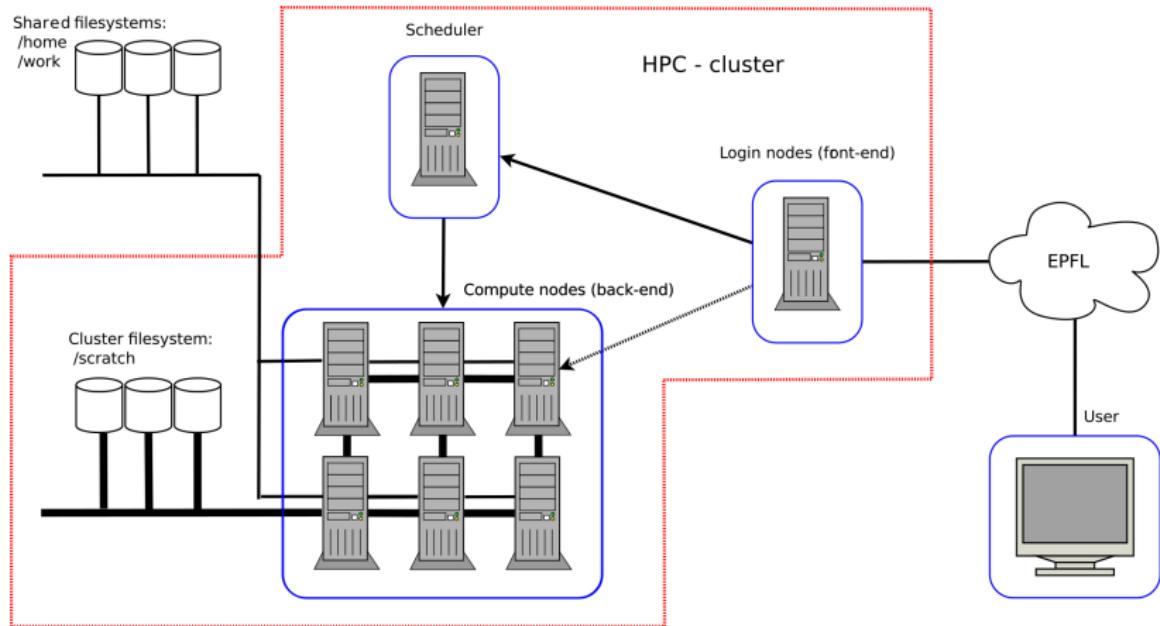
What is a cluster?



What is a cluster?



What is a cluster?



EPFL clusters

Login	Hosts	Cores	RAM	Network
hostname	#	# x GHz	GB	Gbit/s
 fidis.epfl.ch	336	28x2.6	126	56 (IB)
	72	28x2.6	256	
	216	28x2.6	192	100 (IB)
helvetios.epfl.ch	286	36x2.3	192	100 (IB)
izar.epfl.ch	64	40x2.1	192	2x100 (IB)
		+ 2x NVidia V100	7TFlops	
	2	40x2.1	192	100 (IB)
		+ 4x NVidia V100	7.8TFlops	

Remote connection

Connecting to remote machines

First step

- Connect to a remote cluster to get a shell
- SSH: Secure SHell

How to use

- `ssh -l <username> <hostname>`
- `ssh <username>@<hostname>`

Connecting to remote machines

First step

- Connect to a remote cluster to get a shell
- SSH: Secure SHell

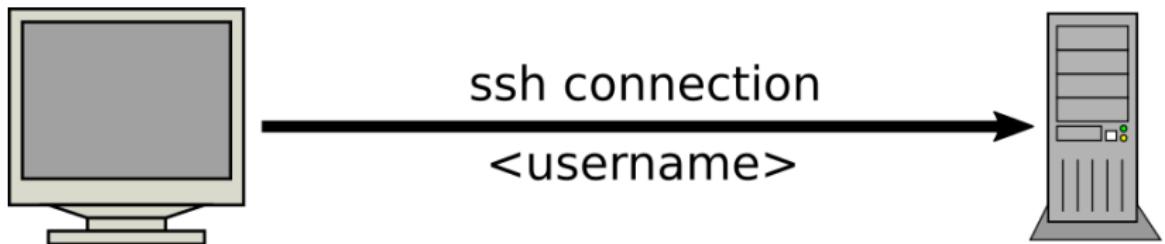
How to use

- `ssh -l <username> <hostname>`
- `ssh <username>@<hostname>`

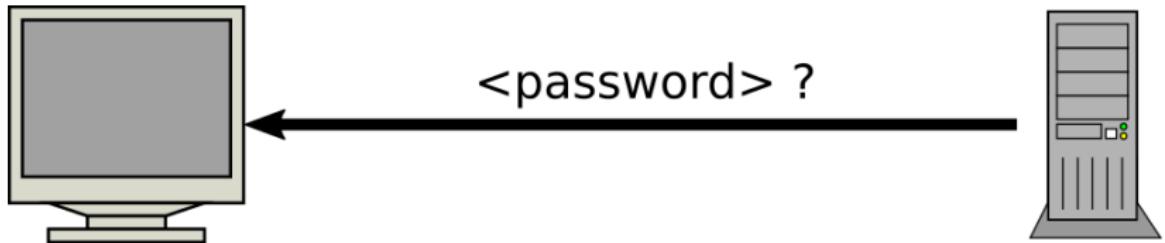
For windows users

Just install git, and use git bash

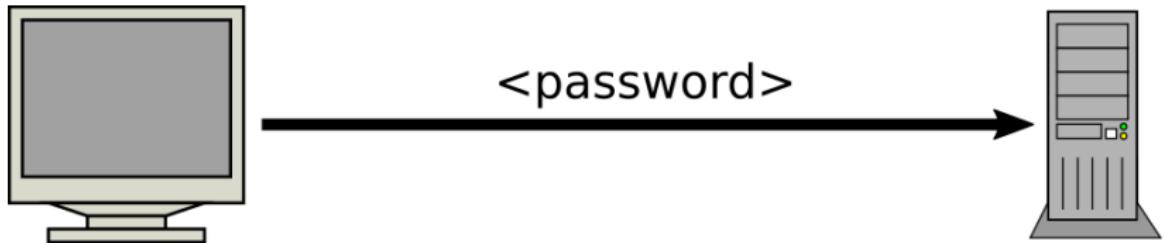
Getting a remote shell



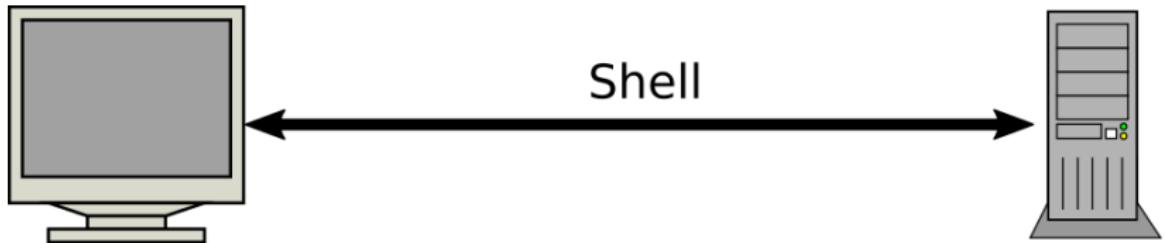
Getting a remote shell



Getting a remote shell



Getting a remote shell



Exercise 1: Simple connection

Front nodes

- **fidis** [CPU (OpenMP/MPI)]
- **izar** [GPU (CUDA)]

Questions:

- Connect to fidis' front node
- Check the different folders `/home` `/scratch`

Transferring files: using scp

scp works fine for Linux/MacOS, it is often pre-installed
On windows use GIT Bash or pscp.exe from Putty

Exercise 2: Using `scp`

scp usage

Send data to remote machine:

```
scp [-r] <local_path> <username>@<remote>:<remote_path>
```

Retrieve data from remote machine:

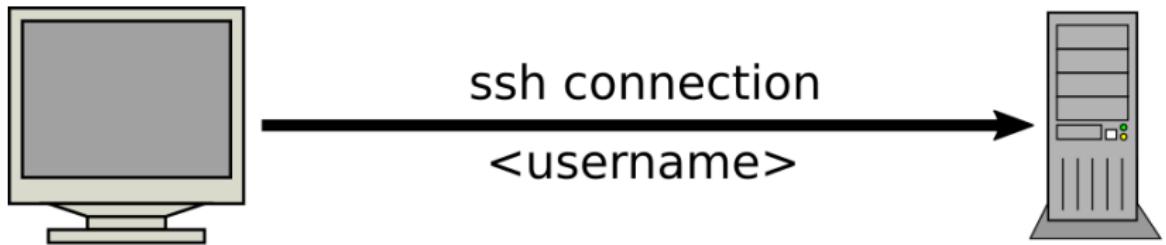
```
scp [-r] <username>@<remote>:<remote_path> <local_path>
```

Note: *It is always easier to “send to” and “receive from” the clusters since they have a fix ip/name*

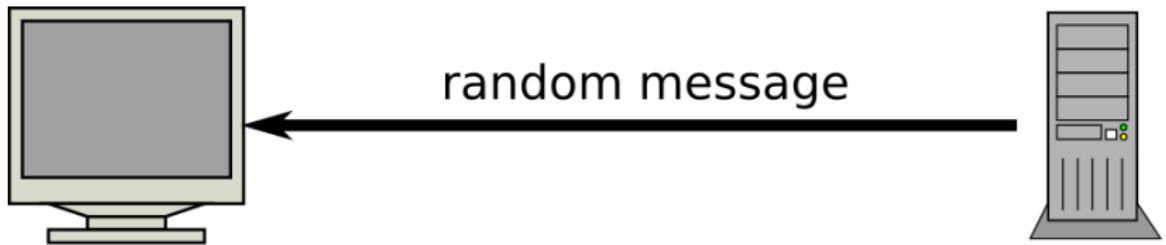
Questions:

- Copy a file from your machine to the cluster.
- Give the name of the file you just copied (leave it as long as you did not get feedbacks for the exercise) **[moodle]**
- Retrieve a file from the cluster onto your machine

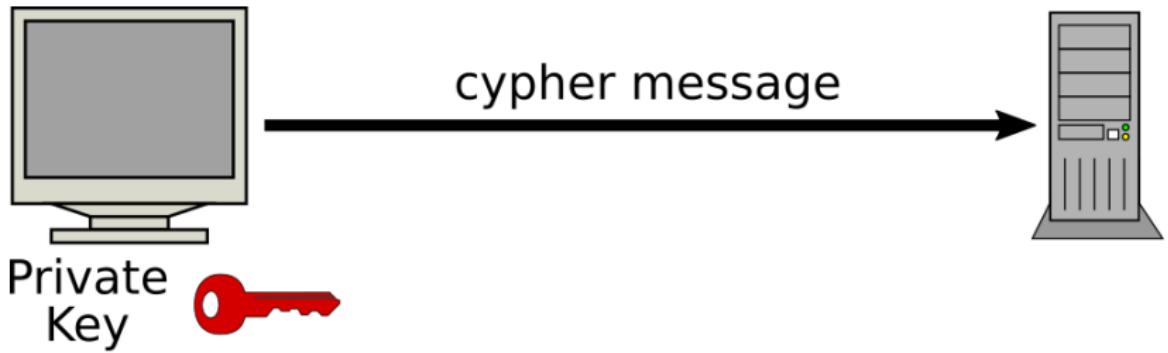
Getting a remote shell using keys



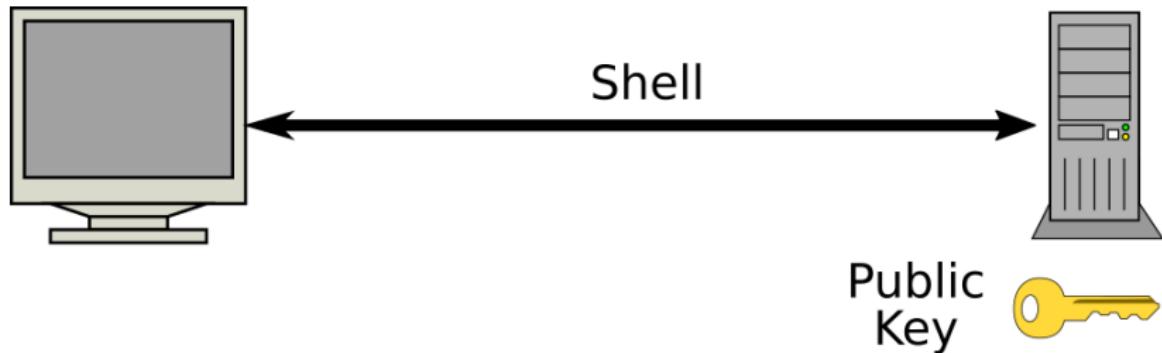
Getting a remote shell using keys



Getting a remote shell using keys



Getting a remote shell using keys



Exercise 3: Connection using keys [optional]

Questions:

- Generate a pair of public/private key using `ssh-keygen`.
- Give us your **public** key **[moodle/optional]**
- Copy the public key `.ssh/id_rsa.pub` using `scp` in a file called `.ssh/authorized_keys` on the remote machine.
- Try to connect (it should not ask your password)

Transferring data using a GUI

sftp

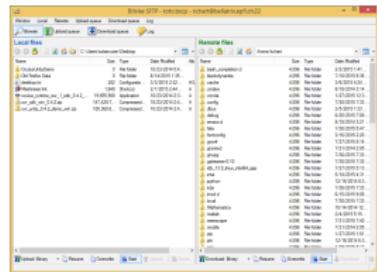
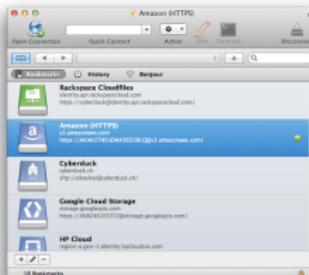
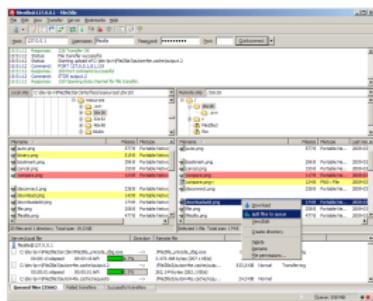
tions of GUIs, for example: Filezilla

<https://filezilla-project.org> ("all" OSes)

Cyberduck <https://cyberduck.io/> (Windows and Mac)

Tunnelier <https://bitvise.com/tunnelier> (Windows)

TUIs also exists like lftp



Using modules

What are modules

- A way to dynamically modify the environment
- They contain configurations to use an application/a library

How to use them

- `module avail` list all possible modules
- `module load <module>` load a module
- `module unload <module>` unload a module
- `module purge` unload all the modules
- `module list` list all loaded modules

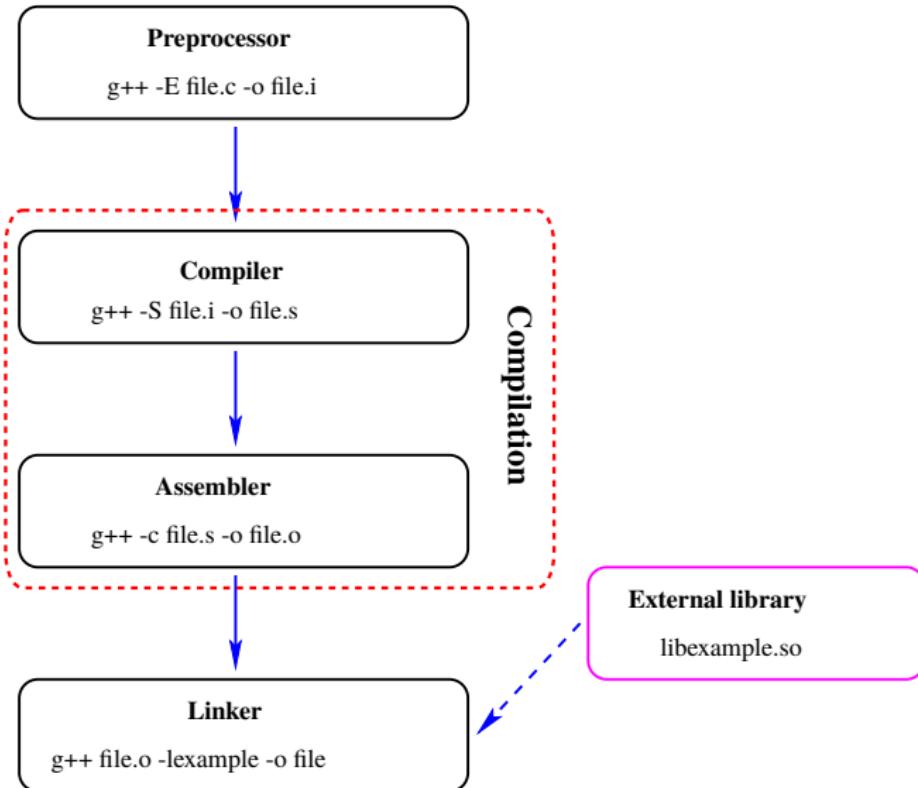
Exercise 4: Module command

Questions:

- List all available modules
- Try `g++ --version`
- Load the module gcc
- Try again `g++ --version`
- Try `icpc --version`
- Load the intel module
- Try again `icpc --version`
- List the currently loaded modules
- Which are the default version of gcc and intel in the modules ? **[moodle]**

Compilation

Compilation stages



Exercise 5: Compilation

Compilation all-in-one

If the compiler is given source code only it will do all the stages at once

Questions:

- Load a compiler module
- Compile the code `hello.c`
- What command did you had to type **[moodle]**
- Try to see the different stages of compilation [Optional]

Separated compilation

- Generate object files: compile

```
g++ -c file1.cc  
g++ -c file2.cc
```

Transforms code file `file?.cc` in object file `file?.o`

- Generate the executable: linking

```
g++ -o hello file1.o file2.o
```

Transforms object file `file?.o` in executable `hello`

Exercise 6: Separated compilation

Questions:

- Enter the folder `hello-separated-files`
- Generate `hello.o`
- Generate `greetings.o`
- Generate `hello` from `hello.o greetings.o`
- What commands did you had to type [**moodle**]

Makefiles

What is makefile

- It is a build automation system
- It is a set of rules on how to produce an executable:
 - what to compile ?
 - how to compile ?
 - what to link ?
 - how to link ?
- It has predefined rules using predefined variables

.CC.O:

```
$(CXX) $(CPPFLAGS) $(CXXFLAGS) -c $<
```

- Add line for linking

\$(EXE): \$(OBJS)

```
$(LD) -o $@ $(LDFLAGS) $(OBJS)
```

- Usage: `make`

Exercise 7: Makefiles

Questions:

- Enter the folder `hello-makefile`
- Read the `Makefile`
- Add the `-Wall -Werror` option to the compilation options
- Compile the code

Generate a library

- Two types of libraries **static** or **shared**

static: archive of object file

shared: library loaded dynamically at execution

- Generate a shared library

Object files have to be compiled with **-fPIC** option

```
g++ -o libgreetings.so -shared greetings.o
```

- Link

```
g++ -L<path to the library> -lgreetings  
-o <executable> <object_1> ... <object_n>
```

Build system

Real life project

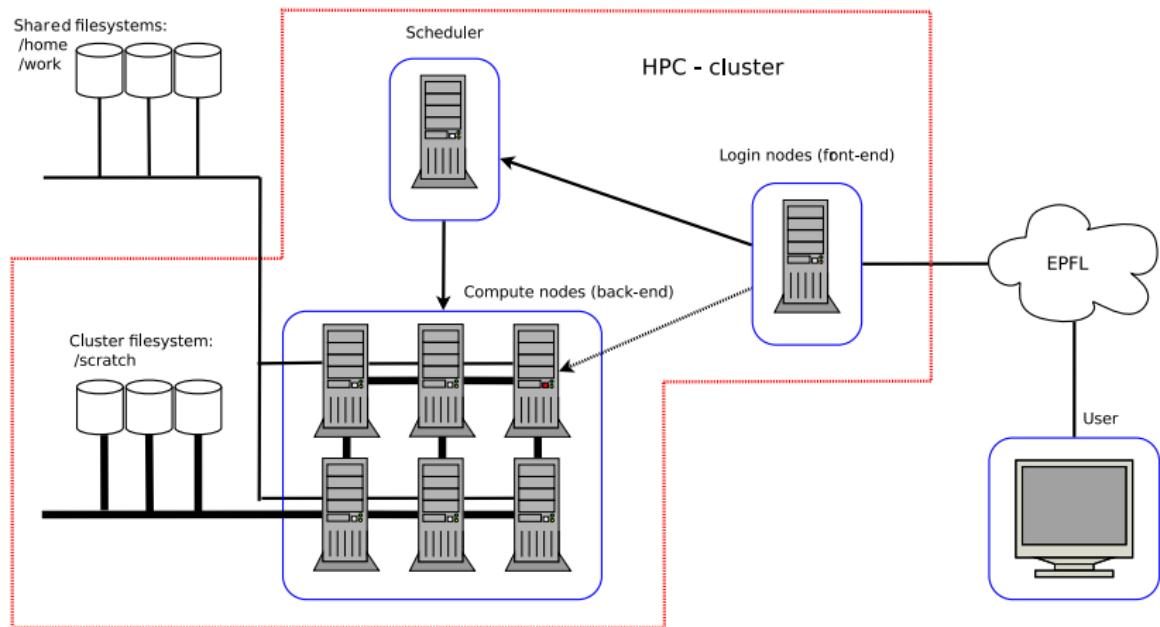
For project a bit more consequent with multiple dependencies to external libraries, usage of a build automation system

- CMake
- Autotools: configure
- SCons

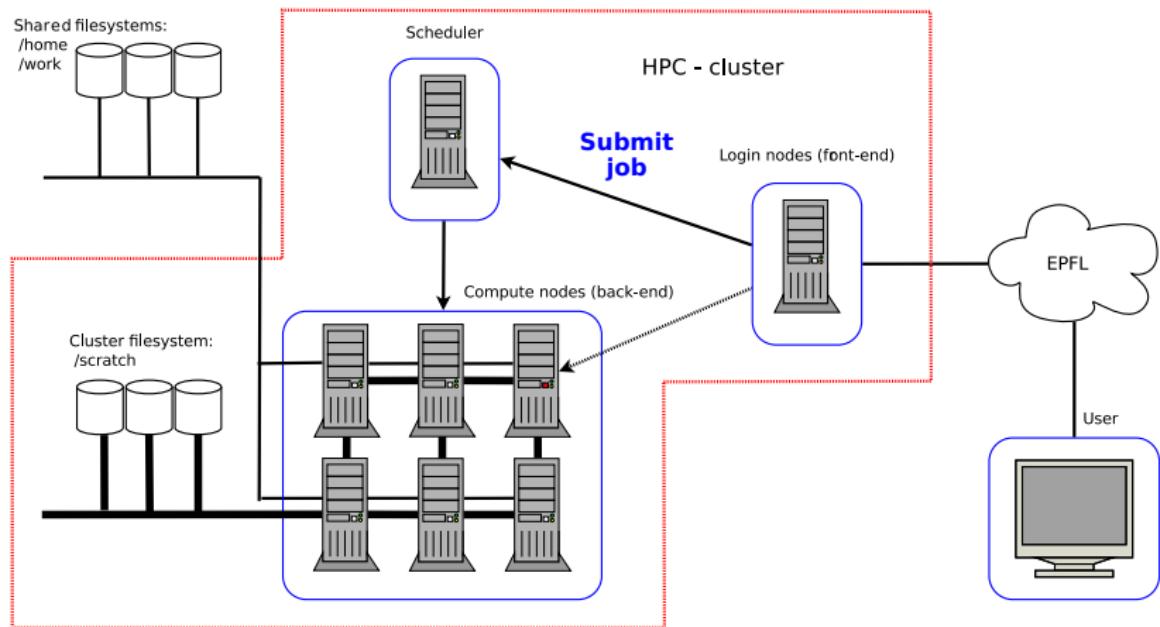
SLURM



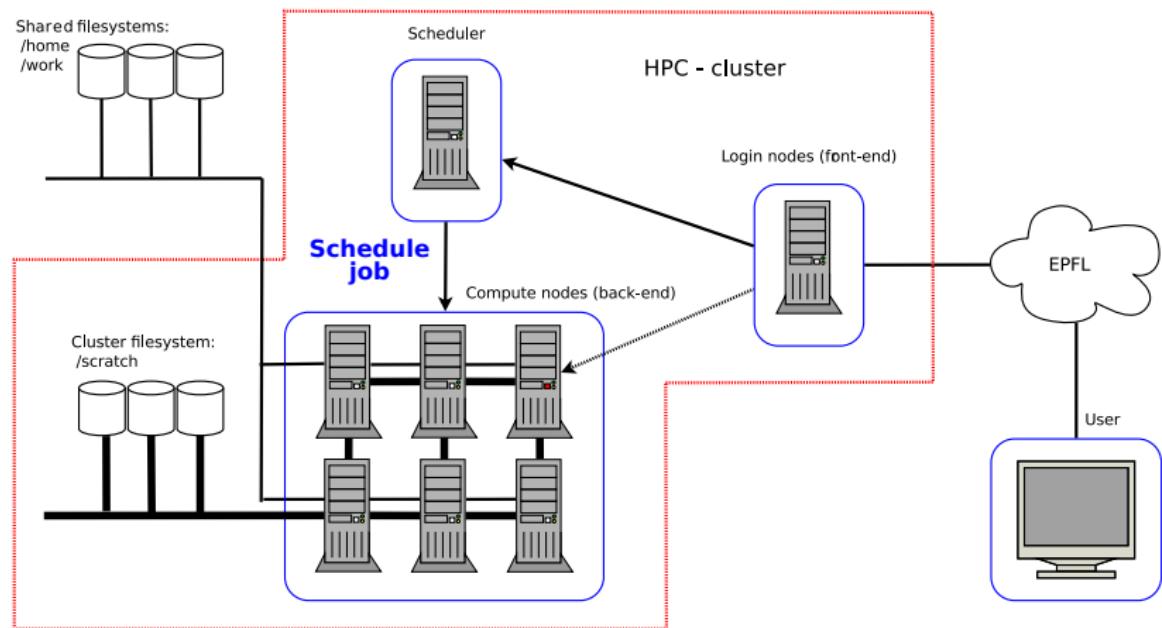
Scheduler



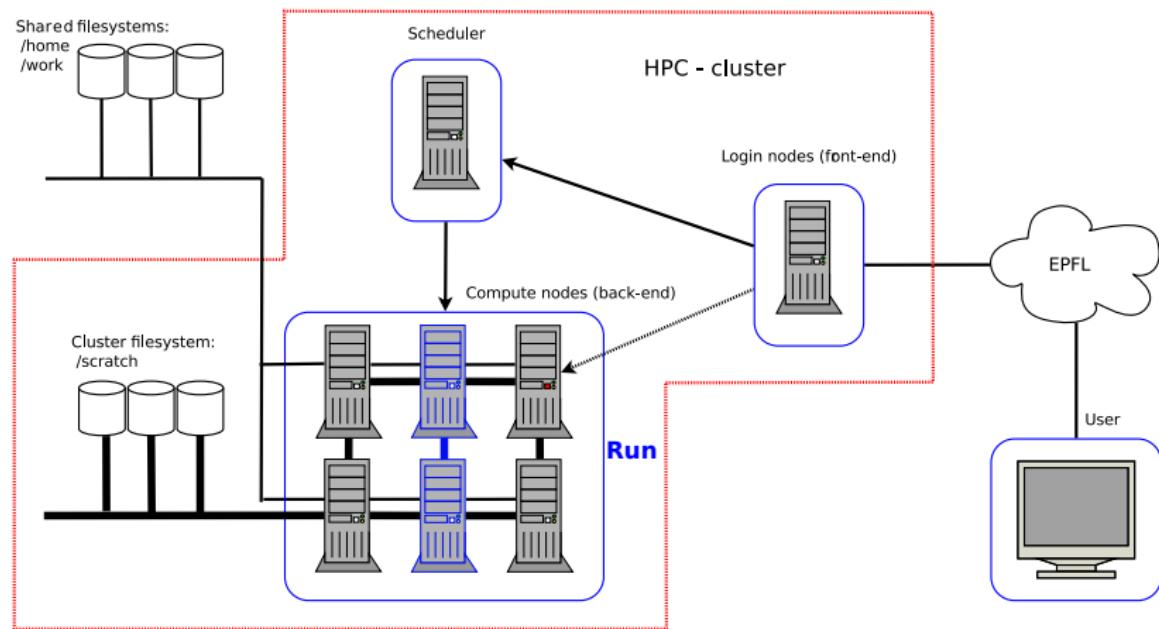
Scheduler



Scheduler



Scheduler



SLURM

What is SLURM

- Simple Linux Utility for Resource Management
- Job scheduler
- Soda in Futurama

Basic commands

- `sbatch` submit a job to the queue
- `salloc` allocates resources
- `squeue` visualize the state of the queue

SLURM: common options

- **-A** `--account` defines which account to use
For this class you are in `phpc2021`
- **-u** `--user` defines the user
mainly useful for `squeue`
- `--reservation` defines which reservation to use
For this class reservations named `phpc2021` are set on both `fidis` and `izar`
- **-p** `--partition` defines the partition to use
List of partition can be get with `sinfo`

SLURM: `sbatch`/`salloc`

- `-N` `--nodes` defines the number of nodes to use
- `-t` `--time` defines the maximum wall time
- `-n` `--tasks` number of tasks, in the MPI sense
- `-c` `--cpus-per-task` number of cpus per process
- `--ntasks-per-node` number of tasks per node
- `--mem` defines the quantity of memory per node requested

`sbatch` job

- A job is simply a shell script
- `sbatch` option can be given as comment `#SBATCH` at the beginning of the script

Exercise 8: SLURM: first commands

Questions:

- Check the queue state
- Use `salloc` to allocate one node in the reservation `phpc2021`
- What command did you had to type? **[moodle]**
- Try `srun hostname`, we will see this command more in detail in later exercise sessions
- Exit the allocation to note block resources: `exit` or `Ctrl-d`

Exercise 9: SLURM: command sbatch

Questions:

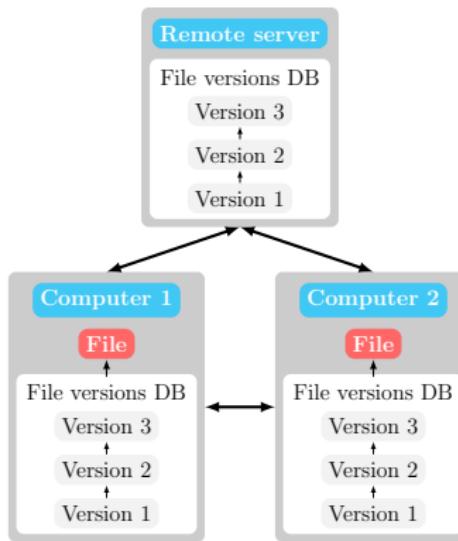
- Write a script that runs the hello world code
- Try your script.
note: *in general you should not try your codes on the front node there is a debug partition for that*
- Submit your script with **sbatch** using the reservation
- Try **squeue -u <username>**
- A file named `slurm-<jobid>.out` should have been created, check its content
- Add the reservation as an option directly in the script
- Submit it again
- Provide the path to your script and do not delete it or change it until you get a feedback **[moodle]**



“Versioning”: with Git

Git: *the stupid content tracker*

- Distributed revision control
- Originally developed by Linus Torvald
- Named after the *egotistical bastard* Linus



git clone



LOCAL SERVER
.GIT DIRECTORY
STAGING AREA

```
$ git clone <uri repo.git>
```



WORKING DIRECTORY

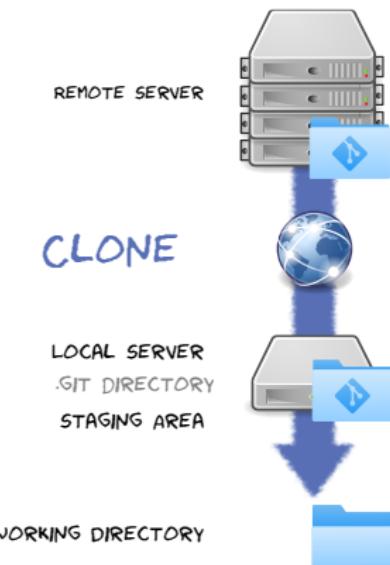
git clone



```
$ git clone <uri repo.git>
Cloning into '<repo>'...
remote: Counting objects: 6940, done.
remote: Total 6940 (delta 0), reused ...
Receiving objects: 100% (6940/6940), ...
Resolving deltas: 100% (3286/3286), done.
```

WORKING DIRECTORY

git clone



```
$ git clone <uri repo.git>
Cloning into '<repo>'...
remote: Counting objects: 6940, done.
remote: Total 6940 (delta 0), reused ...
Receiving objects: 100% (6940/6940), ...
Resolving deltas: 100% (3286/3286), done.
```

git status is your friend

```
$ git status
```

git status is your friend

```
$ git status
```

```
On branch master
```

```
Your branch is up-to-date with 'origin/master'.
```

```
nothing to commit, working tree clean
```

Lets add a file: staging/commit



WORKING DIRECTORY

Lets add a file: staging/commit

```
$ git status
```

```
On branch master
```

```
Your branch is up-to-date with 'origin/master'.
```

Untracked files:

(use "git add <file>..." to include in what will be committed)

my_code.py

nothing added to commit but untracked files present

Lets add a file: staging/commit



Lets add a file: staging/commit

```
$ git status
On branch master
Your branch is up-to-date with 'origin/master'.

Changes to be committed:
  (use "git reset HEAD <file>..." to unstage)

    new file:   my_code.py
```

Lets add a file: staging/commit



COMMIT



LOCAL SERVER
.GIT DIRECTORY
STAGING AREA



```
$ git commit -m <message>
```

WORKING DIRECTORY



Lets add a file: staging/commit

```
$ git status
```

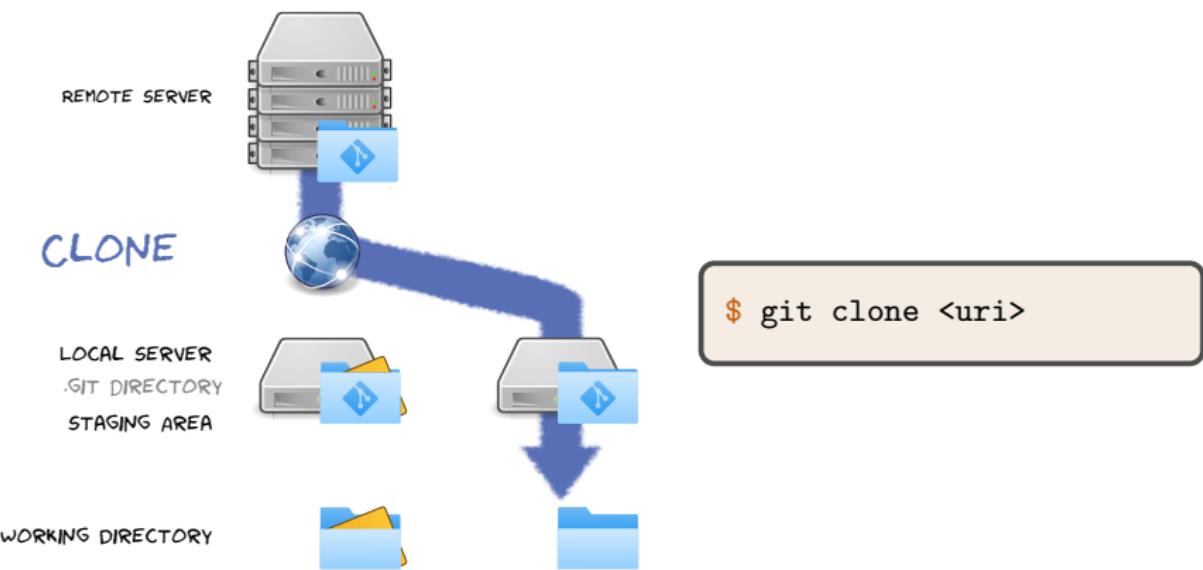
On branch master

Your branch is ahead of 'origin/master' by 1 commit.

(use "git push" to publish your local commits)

nothing to commit, working tree clean

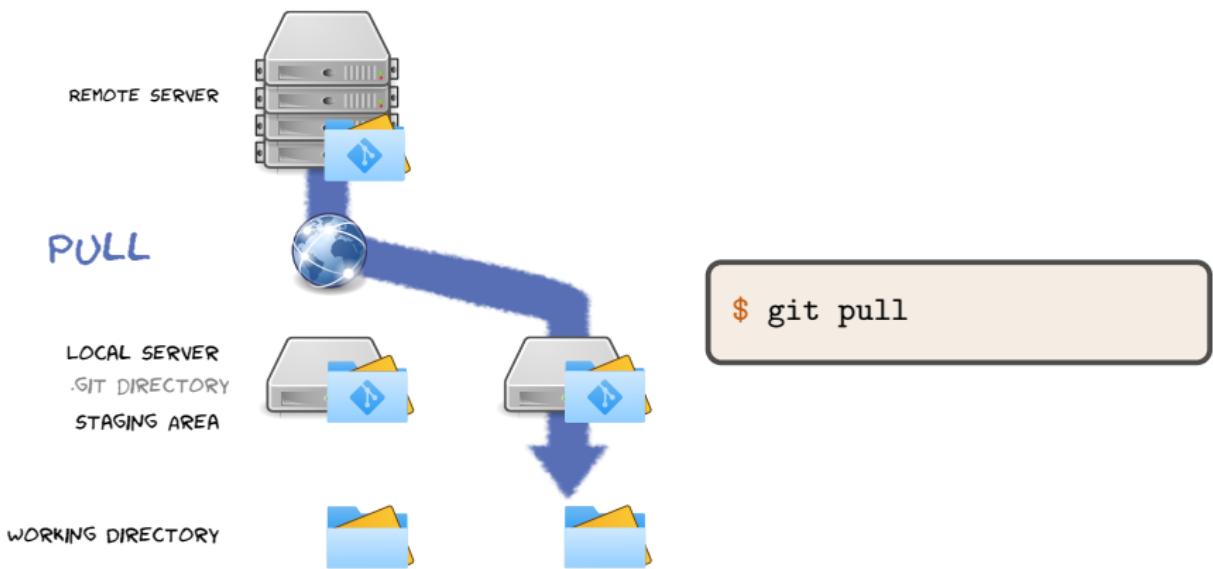
Synchronizing with the remote server



Synchronizing with the remote server



Synchronizing with the remote server



Exercise 10: First step with Git

Questions:

- If you do not have git installed, get it from <https://git-scm.com/downloads> or from your package manager
- Go on <https://c4science.ch/> and login with your EPFL account (Login for Swiss Universities).
- Once connected go on the setting page (the wrench on the top right corner)
- In the **Authentication > SSH Public Keys** menu set your public ssh key. This key will be used to connect to the git server through **ssh**.

Exercise 11: First step with Git

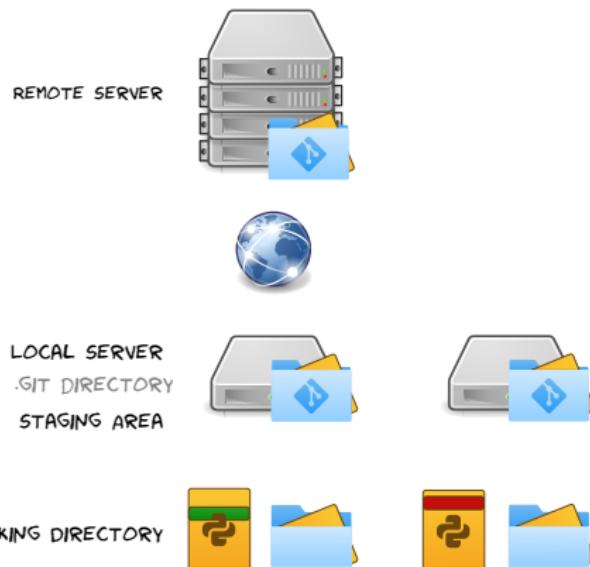
Questions:

- Now you should be able to clone a repository
Either create a repository or clone

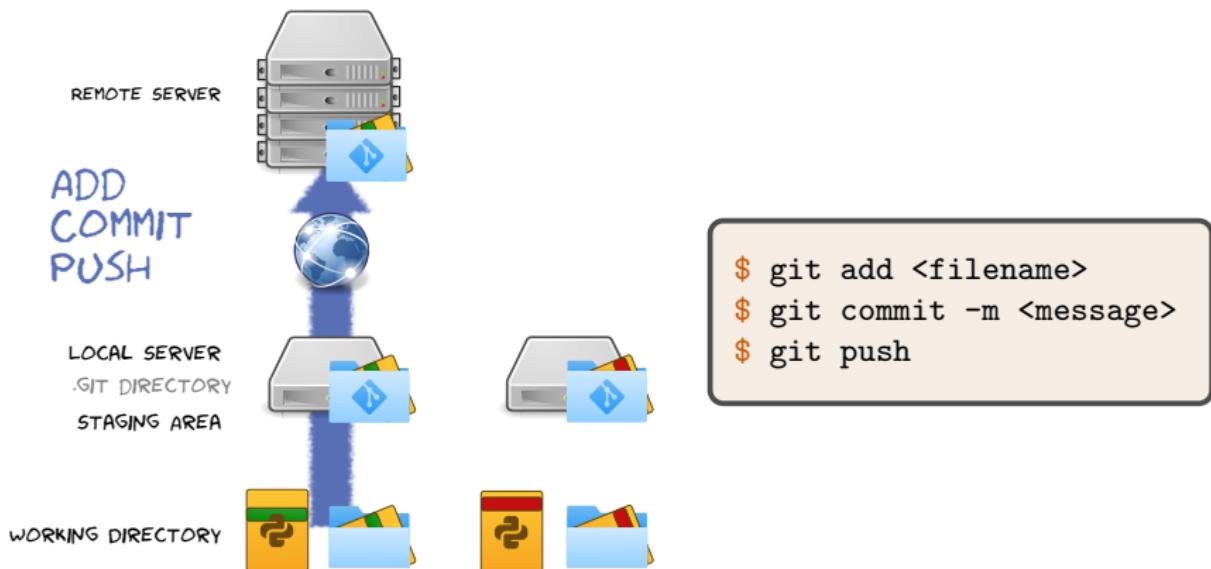
```
ssh://git@c4science.ch/source/scitas-test-repo.git
```

- Create a file, use a filename that will not clash with the others
- Check the state of your working copy
- Add the file to the repository
- Commit your modifications
- Clone the same repository in a different folder
- Pull the potential modifications from the server
- Push your changes to the server
- Give use your c4science username and the name of the file you added **[moodle]**

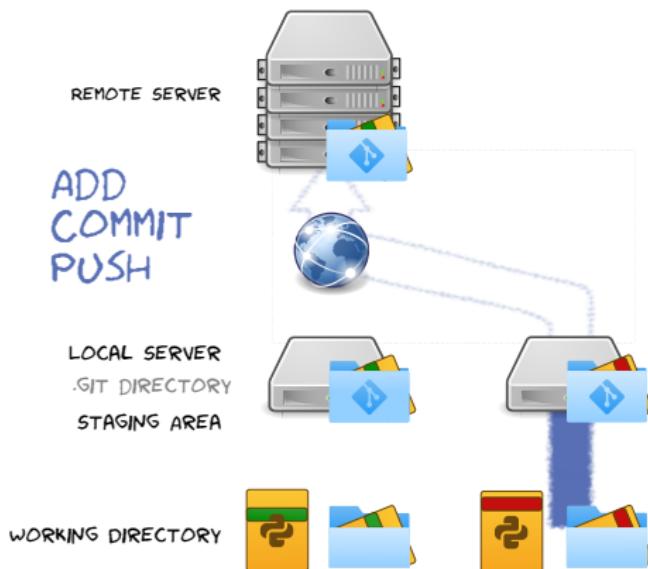
Collaborative work with potential problems



Collaborative work with potential problems



Collaborative work with potential problems

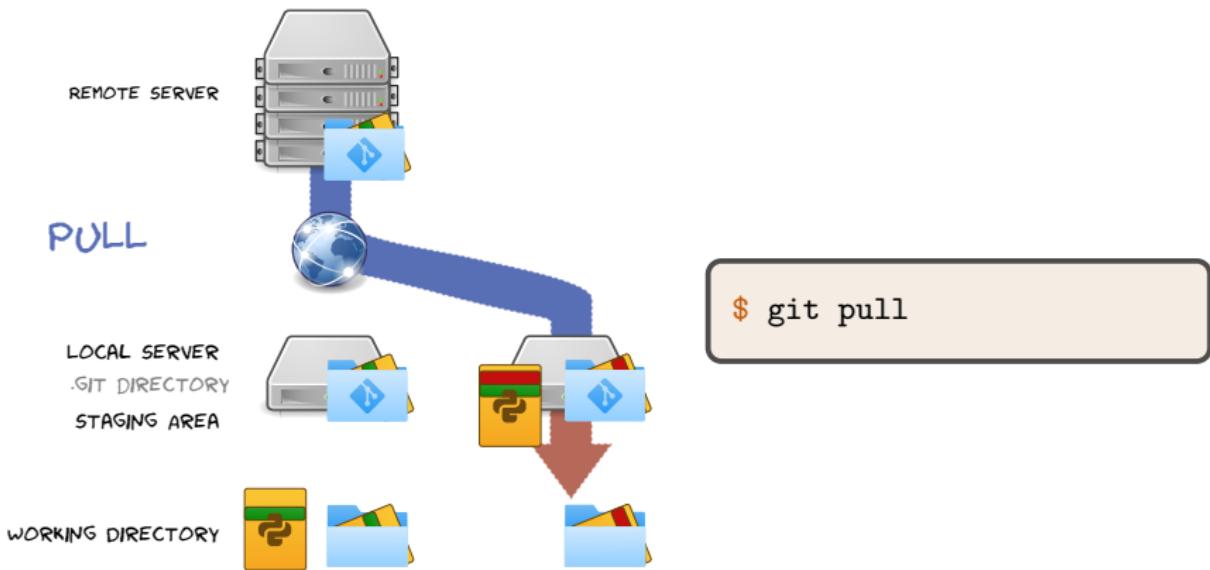


```
$ git add <filename>
$ git commit -m <message>
$ git push
```

Collaborative work with potential problems

```
$ git push
To <repo>
 ! [rejected]           master -> master (fetch first)
error: failed to push some refs to '<repo>'
hint: ...
```

Collaborative work with potential problems



Collaborative work with potential problems

```
$ git pull
remote: Counting objects: 3, done.
remote: Total 3 (delta 0), reused 0 (delta 0)
Unpacking objects: 100% (3/3), done.
From <repo>
  fe22d81..0bcfb99  master      -> origin/master
Auto-merging my_code.py
CONFLICT (content): Merge conflict in my_code.py
Automatic merge failed; fix conflicts and then commit the result.
```

Collaborative work with potential problems

```
$ git status
```

On branch master

Your branch and 'origin/master' have diverged,
and have 1 and 1 different commits each, respectively.

(use "git pull" to merge the remote branch into yours)

You have unmerged paths.

(fix conflicts and run "git commit")

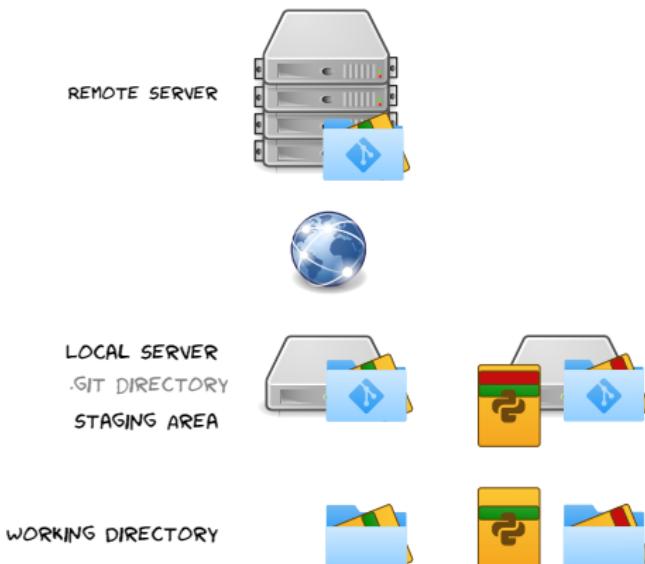
(use "git merge --abort" to abort the merge)

Unmerged paths:

(use "git add <file>..." to mark resolution)

both modified: my_code.py

Collaborative work with potential problems



Correct the conflict:

my_file.py

<<<<<<<

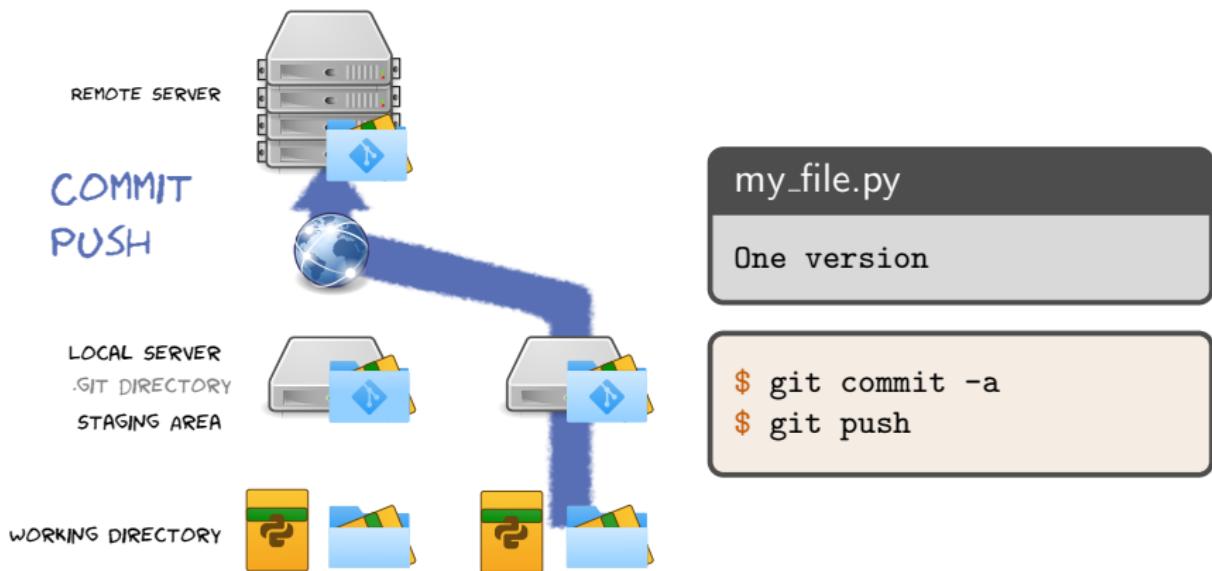
One version

=====

Other version

>>>>>>

Collaborative work with potential problems



Exercise 12: Generate and solve conflicts

Questions:

- Modify the file created in the previous exercise in both clones
- Commit this both modifications
- Pull and push in one of the clone
- Pull in the second clone, You should get a conflict

<<<<<<<

One version

=====

Other version

>>>>>>

- Check the local status
- Correct the conflict and commit using `git commit -a`
- Push the modifications

Introduction to branches



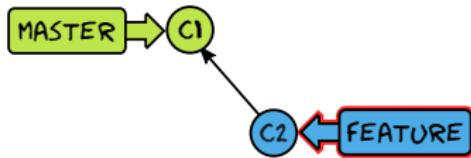
```
$ git clone <uri repo.git>
```

Introduction to branches



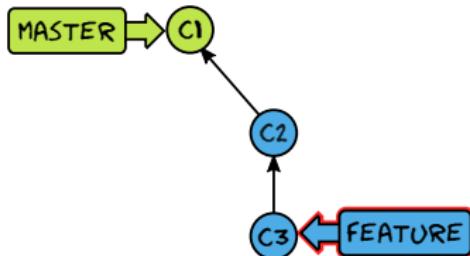
```
$ git checkout -b feature
```

Introduction to branches



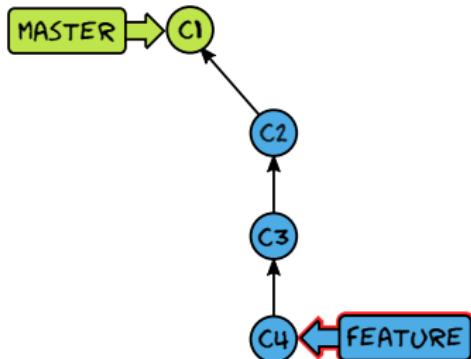
```
$ git commit -m <message>
```

Introduction to branches



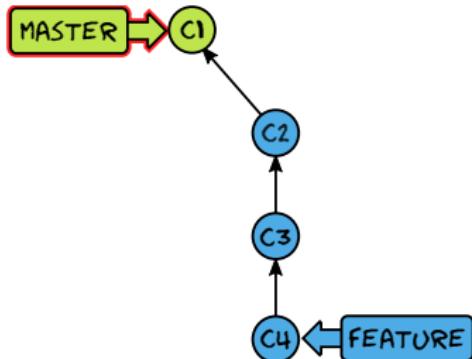
```
$ git commit -m <message>
```

Introduction to branches



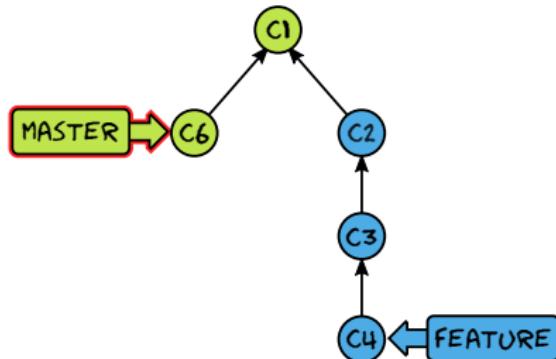
```
$ git commit -m <message>
```

Introduction to branches



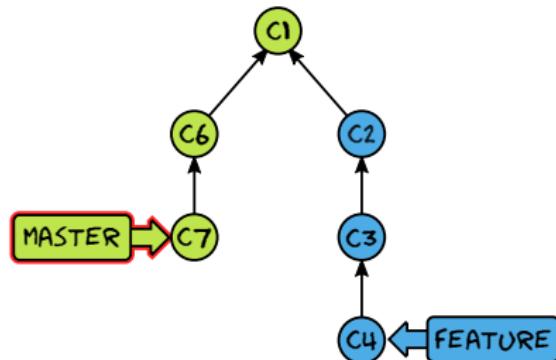
```
$ git checkout master
```

Introduction to branches



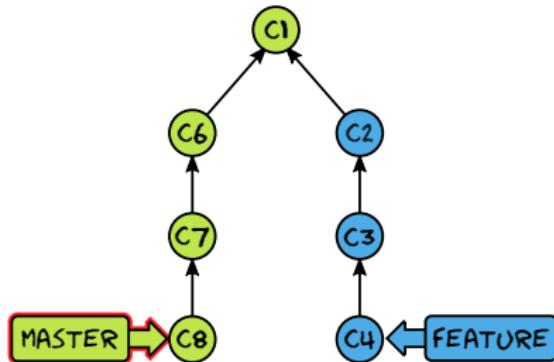
```
$ git commit -m <message>
```

Introduction to branches



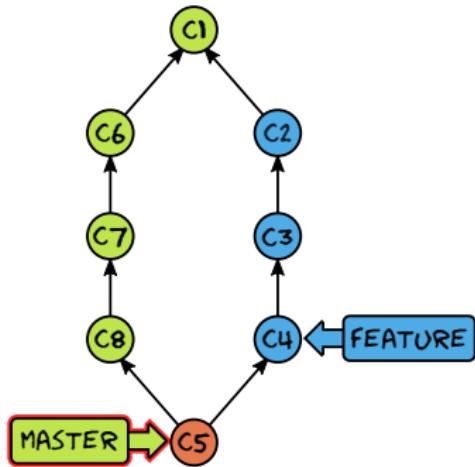
```
$ git commit -m <message>
```

Introduction to branches



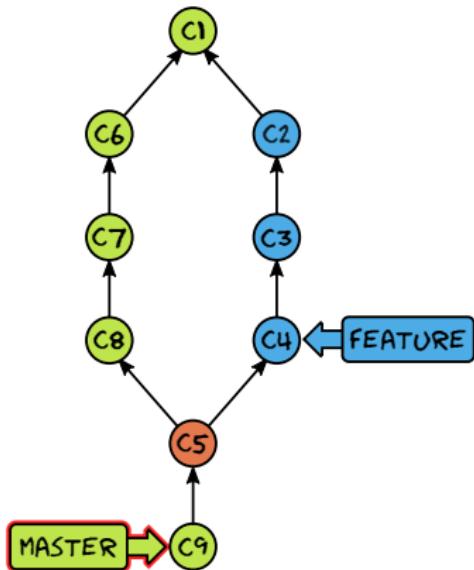
```
$ git commit -m <message>
```

Introduction to branches



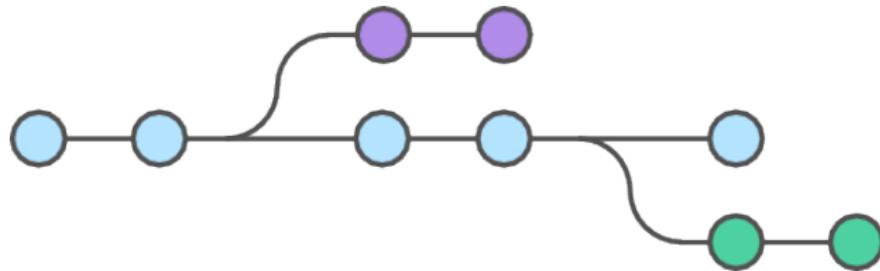
```
$ git merge feature
```

Introduction to branches



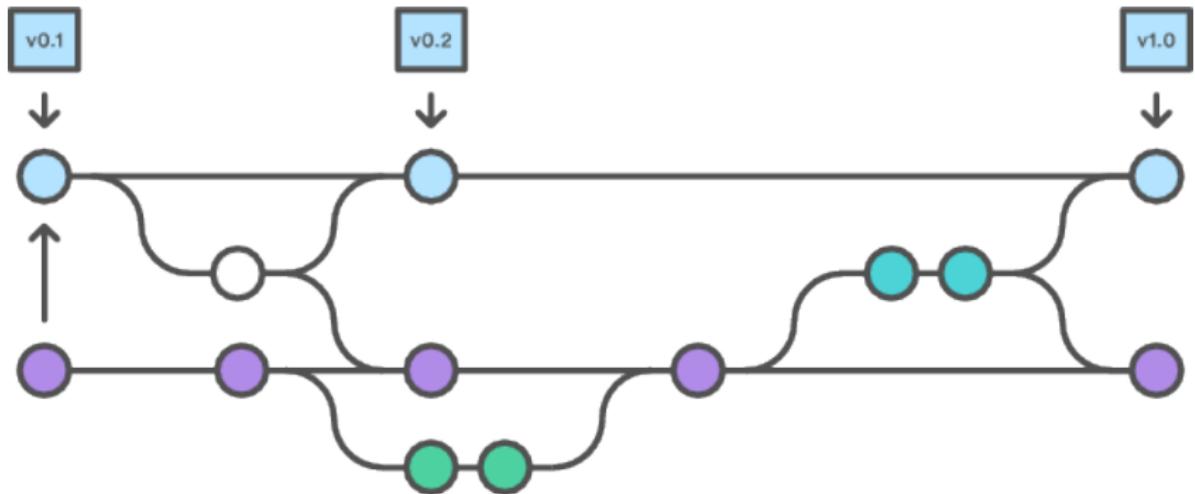
```
$ git commit -m <message>
```

Feature branch

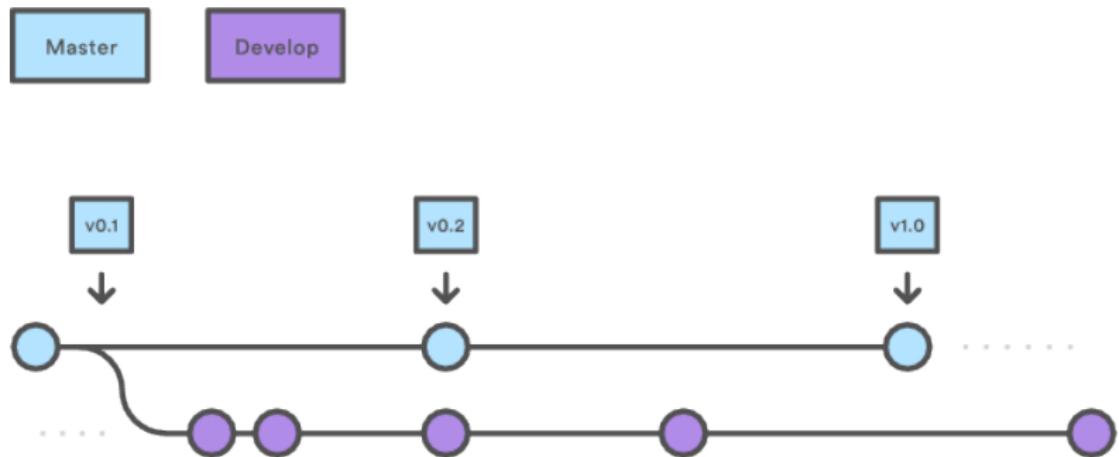


Workflow: gitflow

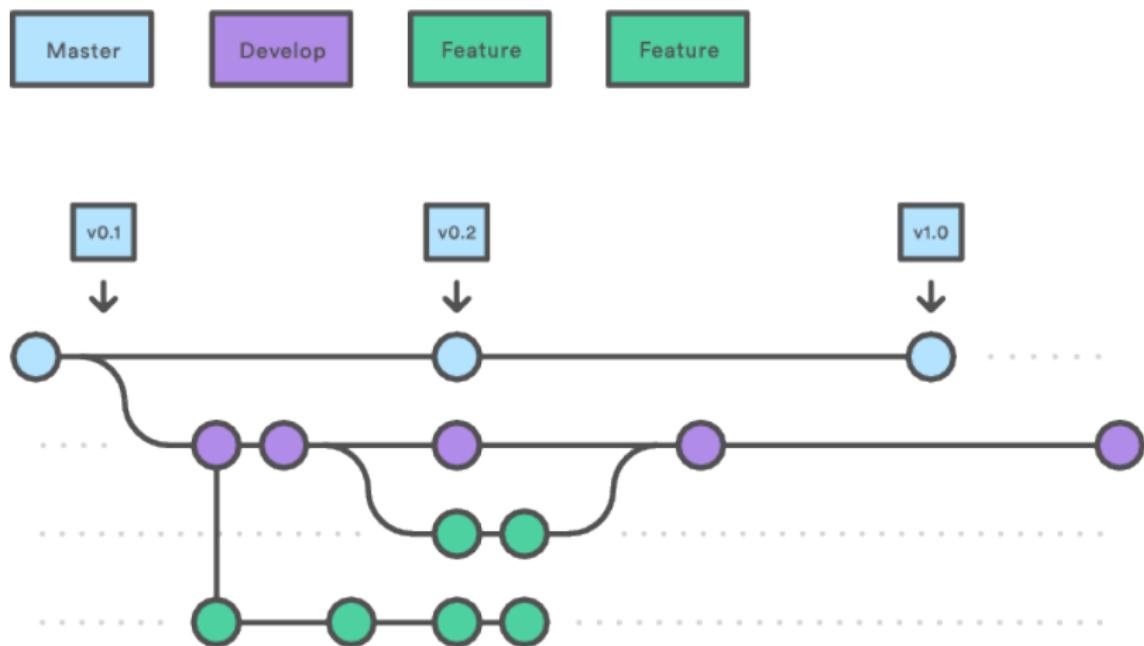
Gitflow



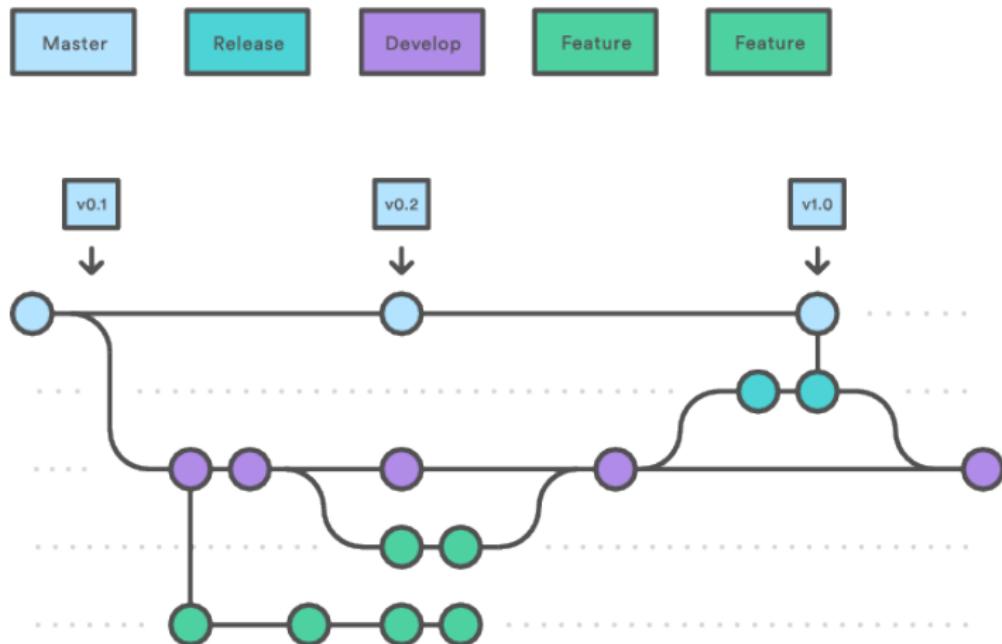
Workflow: gitflow



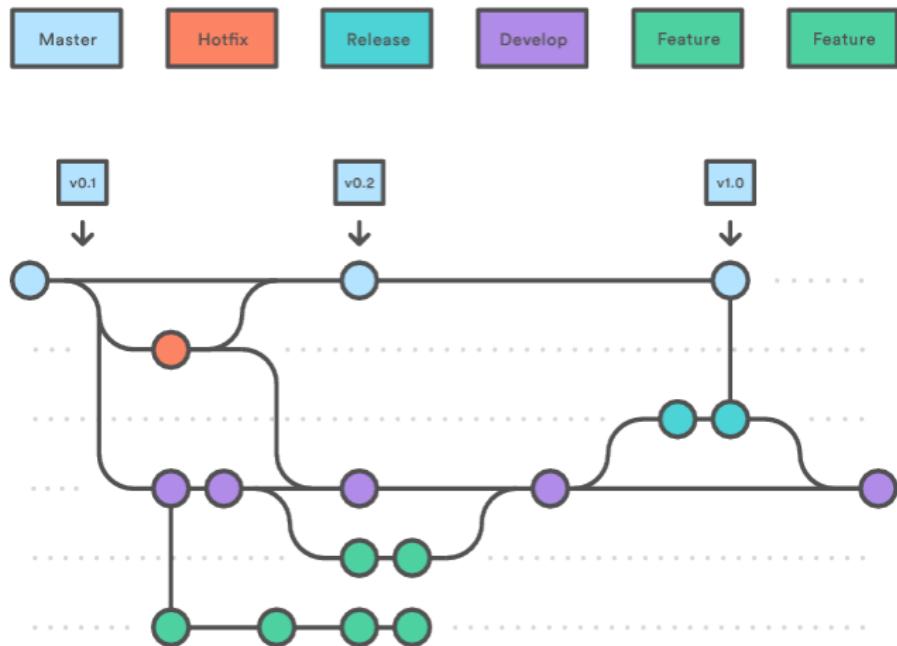
Workflow: gitflow



Workflow: gitflow



Workflow: gitflow



Exercise 13: Branches/merges

Questions:

- Create a branch with the name of your choice
- Modify a file and commit the changes
- Checkout the `master` branch
- Modify a file and commit the changes
- Merge the branch previously created in the `master` branch
- List all branches
- Print the logs of the different modifications
- Give us the name of the branch you created **[moodle]**

Sources and extra infos

Sources

- Wikipedia
- <http://git-scm.com>
- Manpages: rsync, git
- <https://www.atlassian.com/git/>
- [http://nvie.com/posts/
a-successful-git-branching-model/](http://nvie.com/posts/a-successful-git-branching-model/)

Learn more

- Git with a game: <http://learngitbranching.js.org/>