

A Short Introduction to OpenMP

- **Open Multi-Processing**
 - API for shared memory multi-processing in Fortran, C and C++
- **Simpler than MPI, and simpler than pthread**
 - Coarse and fine grain parallelisation
 - Fork-join model
 - Since version 3.0, supports arbitrary tasks
- **Compiler directives and environment variables**
 - Requires compiler support
 - Intel compiler: OpenMP v2.5, GCC 4.2: v2.5, GCC 4.4: v3.0, GOMP extension for earlier GCC versions, Microsoft Visual Studio 2005: v2.0, Sun Studio: v2.5, Pathscale: v2.5

```
#include <omp.h>
#include <stdio.h>

int main(int argc, char *argv[])
{
    #pragma omp parallel
        printf("Hello from thread %d out of %d",
               omp_get_thread_num(), omp_get_num_threads());
    printf("Hello again from thread %d",
           omp_get_thread_num());
    return 0;
}
```

Compiling and executing

```
%> gcc -fopenmp -Wall hello.c
%> icc -openmp -Wall hello.c
%> export OMP_NUM_THREADS=4
%> ./a.out
Hello from thread 0 out of 4
Hello from thread 1 out of 4
Hello from thread 2 out of 4
Hello from thread 3 out of 4
Hello again from thread 0
```

- The number of threads can also be set within the source code using the *omp_set_num_threads* function

All the examples assume that 4 threads are available

- **Synchronization at beginning and end of parallel regions**
 - Can be refined via **synchronization directives**
- **By default, threads in parallel regions share all variables**
 - Visibility can be refined via **data placement directives**

Parallel sections

```
int main(int argc, char *argv[])
{
    #pragma omp parallel sections
    {
        #pragma omp section
            printf("Hello from thread %d out of %d",
                omp_get_thread_num(), omp_get_num_threads());
        #pragma omp section
            printf("Hello again from thread %d",
                omp_get_thread_num());
    }
    return 0;
}
```

Output:

```
Hello from thread 0 out of 4
Hello again from thread 1
```

Parallel for loops

```
#pragma omp parallel
{
#pragma omp for
    for (int i=0; i<size; ++i)
    {
        do stuff in parallel
    }
}
```

- **Loop index is private within each thread**
- **Fortran**
 - Parallel DO loops
 - Parallel workshare (exists for C/C++ but not supported by current compilers)

Parallel for loops: scheduling

- **schedule** directive describes distribution of iterations onto threads
- **RUNTIME**
 - Default strategy, uses strategy defined in OMP_SCHEDULE environment variable, or static by default
- **STATIC**
 - Iterations are pre-assigned to threads before loop execution, iterations are grouped into chunks
 - Chunk size parameter specifies iterations count per block
- **DYNAMIC**
 - Chunks are assigned dynamically to threads
 - Chunk size parameter specifies iterations count per block
- **GUIDED**
 - Block sizes are decreasing according to exponential law
 - Chunk parameter specifies minimum chunk size to be used

- **Good load balancing depends on**

- The algorithm
- The scheduling strategy
- The specified chunk size

```
#pragma omp parallel for schedule(static, 200)
  for (int i=0; i<size; ++i)
  {
    if (i<20) do stuff else do nothing
  }
```

- Here: bad load balance for chunk size of 200, much better if chunk size is 5
- For **dynamic** strategy: small chunk → more overheads, large chunks → less chunks to distribute

Data placement directives

- Except for loop indices, variables are shared by default
- Variables on the stack can be made private to be replicated by each thread
 - Only variables on the stack can be made private
- **private**
 - Variable is undefined at beginning of parallel region
- **firstprivate**
 - Variables takes last value outside parallel region
- **lastprivate**
 - Only for parallel for loops, after exiting the loop, the variable takes the value reached in the last iteration
- **shared**
 - Default behavior

Data placement: shared

```
int i; int mynum = 10; int a = 10;

printf("a=%d before parallel region\n", a);
#pragma omp parallel for
for (i=0; i<4; ++i) {
    mynum = omp_get_thread_num();
    a += mynum;
    printf("a=%d in thread %d\n", a, mynum);
}
printf("a=%d after parallel region\n", a);
```

Output:

```
a=10 before parallel region
a=10 in thread 0
a=11 in thread 1
a=13 in thread 2
a=16 in thread 3
a=16 after parallel region
```

Data placement: private

```
int i; int mynum = 10; int a = 10;
printf("a=%d before parallel region\n", a);
#pragma omp parallel for private(a)
for (i=0; i<4; ++i) {
    mynum = omp_get_thread_num();
    a += mynum;
    if (mynum==0) sleep(1); /* thread 0 sleeps 1 second */
    printf("a=%d in thread %d\n", a, mynum);
}
printf("a=%d after parallel region\n", a);
```

Output:

a=10 before parallel region

a=10 in thread 3

a=8 in thread 1

a=9 in thread 2

a=6322816 in thread 3

a=10 after parallel region

mynum is indeed shared



Data placement: firstprivate

```
int i; int mynum = 10; int a = 10;
printf("a=%d before parallel region\n", a);
#pragma omp parallel for firstprivate(a, mynum)
for (i=0; i<4; ++i) {
    mynum = omp_get_thread_num();
    a += mynum;
    printf("a=%d in thread %d\n", a, mynum);
}
printf("a=%d after parallel region\n", a);
```

Output:

```
a=10 before parallel region
a=13 in thread 3
a=12 in thread 2
a=11 in thread 1
a=10 in thread 0
a=10 after parallel region
```

Data placement: lastprivate

```
int i; int mynum = 10; int a = 10;
printf("a=%d before parallel region\n", a);
#pragma omp parallel for lastprivate(a) private(mynum)
for (i=0; i<4; ++i) {
    mynum = omp_get_thread_num();
    a += mynum;
    printf("a=%d in thread %d\n", a, mynum);
}
printf("a=%d after parallel region\n", a);
```

Output:

```
a=10 before parallel region
a=6363522 in thread 2
a=6380803 in thread 3
a=6317952 in thread 0
a=6346113 in thread 1
a=6380803 after parallel region
```

Synchronization directives

- **Operations**
 - ATOMIC
 - BARRIER
 - REDUCTION
- **Blocs**
 - CRITICAL
 - SINGLE
- **Buffer flush**
 - FLUSH
- **Remove synchronization**
 - NOWAIT
- **Remove synchronization**

- Accesses to a shared variable must be protected

```
#pragma omp parallel for
for (int i=0; i<N; ++i) {
    for (int j=0; j<N; ++j) {
        #pragma omp atomic
        sum += A[i][j] ;
    }
}
```

- Equivalent to using *pthread* locks, but much simpler

- All-to-one operation in a parallel for


```
#pragma omp parallel for reduction(+:_sum)
for (int i=0; i<N; ++i) {
    for (int j=0; j<N; ++j) {
        _sum += A[i][j]
    }
}
```

- The OpenMP runtime sums the values of each thread at the end of the loop

Single directive

- Only first incoming thread executes a **single** region
- The other threads wait, unless **nowait** is specified


```
#pragma omp parallel
{
    ...
    #pragma omp single
    {
        ...
    }
    ...
}
```



Critical directive

- At a given time, only one thread can be in a **critical** region

```
#pragma omp parallel
{
    ...
    #pragma omp critical
    {
        ...
    }
    ...
}
```



- **OpenMP 2.0**

<http://www.openmp.org/mp-documents/cspec20.pdf>

- **OpenMP 2.5**

<http://www.openmp.org/mp-documents/spec25.pdf>

- **OpenMP 3.0**

<http://www.openmp.org/mp-documents/spec30.pdf>